

# MapReduce

# MapReduce?

Reason	Explanation
1. Scalability	Efficiently processes terabytes/petabytes by distributing the workload.
2. Parallelism	Exploits <a href="#">data parallelism</a> with minimal developer effort.
3. Fault Tolerance	Built-in <a href="#">recovery of failed tasks</a> via task trackers and job schedulers.
4. Simplicity	<a href="#">Developers focus only on writing map() and reduce() logic.</a>
5. Cost Efficiency	Runs on commodity hardware using <a href="#">open-source</a> tools like Hadoop.
6. Data Locality	Moves computation to the data, not vice versa, <a href="#">saving bandwidth.</a>
7. Ecosystem Integration	<a href="#">Integrates well</a> with tools like Hive, Pig, Spark, HBase, etc.

# MapReduce Framework

What is **MapReduce**?

- **Programming model + implementation**
- Developed by Google in 2008

*Google:*

A simple and powerful interface that enables **automatic parallelization and distribution of large-scale computations**, combined with an implementation of this interface that achieves high performance on **large clusters of commodity PCs**.

# History and Motivation

## Google PageRank problem (2003)

- How to rank tens of billions of web pages by their importance
  - ... efficiently in a reasonable amount of time
  - ... when data is scattered across thousands of computers
  - ... data files can be enormous (terabytes or more)
  - ... data files are updated only occasionally (just appended)
  - ... sending the data between compute nodes is expensive
  - ... hardware failures are rule rather than exception
- Centralized index structure was no longer sufficient
- Solution
  - **Google File System** – a distributed file system
  - **MapReduce** – a programming model

# MapReduce Framework

## MapReduce **programming model**

- **Cluster** of commodity personal computers (nodes)
  - Each running a host operating system, mutually interconnected within a network, communication based on IP addresses, ...
- **Data is distributed among the nodes**
- **Tasks executed in parallel across the nodes**

## Classification

- Process interaction: **message passing**
- Problem decomposition: **data parallelism**



# Basic Idea

## **Divide-and-conquer** paradigm

- Breaks down a given problem into simpler sub-problems
- Solutions of the sub-problems are then combined together

## Two core functions

- **Map function**
  - Generates a set of so-called **intermediate key-value pairs**
- **Reduce function**
  - Reduces values associated with a given intermediate key

And that's all!

# Basic Idea

And that's really all!

It means...

- We only need to **implement *Map* and *Reduce* functions**
- **Everything else** such as
  - input data distribution,
  - scheduling of execution tasks,
  - monitoring of computation progress,
  - inter-machine communication,
  - handling of machine failures,
  - ...

**is managed automatically** by the framework!

# Model Description

## Map function

- *Input*: **input key-value pair** = *input record*
- *Output*: **list of intermediate key-value pairs**
  - Usually from a different domain
  - Keys do not have to be unique
  - Duplicate pairs are permitted
- $(key, value) \rightarrow \text{list of } (key, value)$

## Reduce function

- *Input*: **intermediate key + list of (all) values** for this key
- *Output*: **possibly smaller list of values** for this key
  - Usually from the same domain
- $(key, \text{list of values}) \rightarrow (key, \text{list of values})$



# Example: Word Frequency

```
/**
 * Map function
 * @param key    Document identifier
 * @param value  Document contents
 */
map(String key, String value) {
    foreach word w in value: emit(w, 1);
}
```

where,

- *value* is a line of text from the input (e.g., a sentence).
- *for each word w in value*: this loops through each word in the sentence.
- *emit(w, 1)*: for every word w, output a key-value pair where:
  - the *key* is the word itself,
  - the *value* is 1, representing one occurrence of that word.

Example:

Input line (*value*):

"cat dog cat"

Map function output:

("cat", 1)

("dog", 1)

("cat", 1)

} Intermediate key value

These key-value pairs are then sent to the Reduce function, which adds up all the counts for each word.

After the `map()` step emits these pairs, the framework groups all values by key and sends them to `reduce()`.

# Example: Word Frequency

```
/**
 * Reduce function
 * @param key    Particular word
 * @param values List of count values generated for this word
 */
reduce(String key, Iterator values) {
    int result = 0;
    foreach v in values: result += v;
    emit(key, result);
}
```

*values*: A list/collection of integers associated with a specific key (**from the Map step**).

*foreach v in values*: Loop through each value.

*result += v*: Add each value to result.

*emit(key, result)*: final output of that key.

**For example:**

reduce("cat", [1, 1]) → ("cat", 2)

reduce("dog", [1]) → ("dog", 1)

E.g., a word "cat" was emitted like this from the Map step for **another sentence containing cat 3 times**:

("cat", 1)

("cat", 1)

("cat", 1)

The **framework groups these by key** and sends key "cat" and a list of values [1, 1, 1] to Reducer.

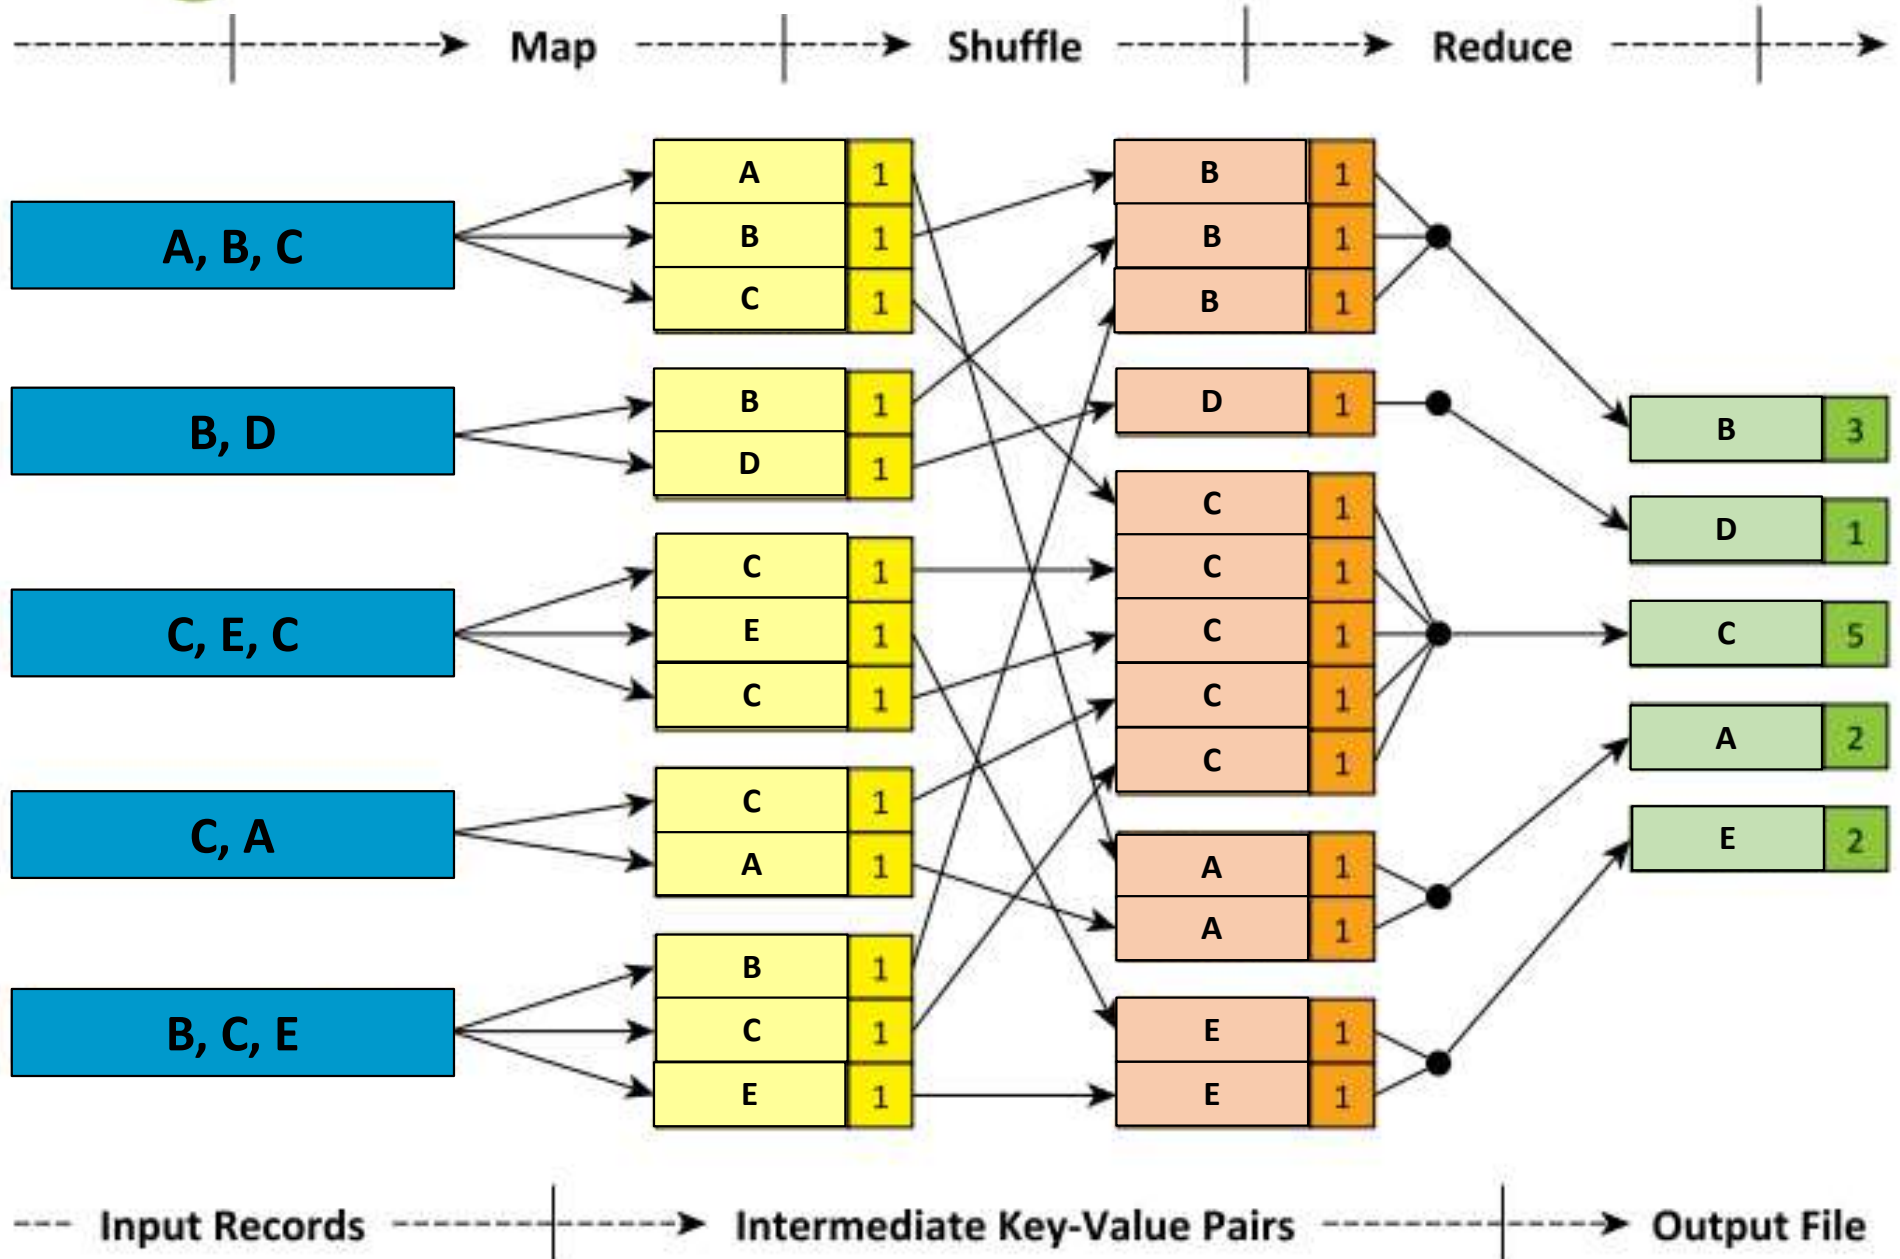
**Reducer Code:**

int result = 0;

foreach v in [1, 1, 1]: result += v;

emit("cat", result); **// Output: ("cat", 3)**

# Logical Phases



# Logical Phases

## Mapping phase

- **Map function** is executed **for each input record**
- Intermediate key-value pairs are emitted

## Shuffling phase

- Intermediate key-value pairs are **grouped and sorted** according to the keys

## Reducing phase

- **Reduce function** is executed **for each intermediate key**
- Output key-value pairs are generated


# Cluster Architecture

## Master-slave architecture


- Two types of nodes, each with two basic roles
- **Master**
  - **Manages the execution of MapReduce jobs**
    - Schedules individual Map / Reduce tasks to idle workers
    - ...
  - **Maintains metadata about input / output files**
    - These are stored in the underlying distributed file system
- **Slaves (workers)**
  - **Physically store the actual data contents of files**
    - Files are divided into smaller parts called splits
    - Each split is stored by one / or even more particular workers
  - **Accept and execute assigned Map / Reduce tasks**




# Cluster Architecture



# MapReduce Job Submission



# Execution Schema



# MapReduce Job Submission

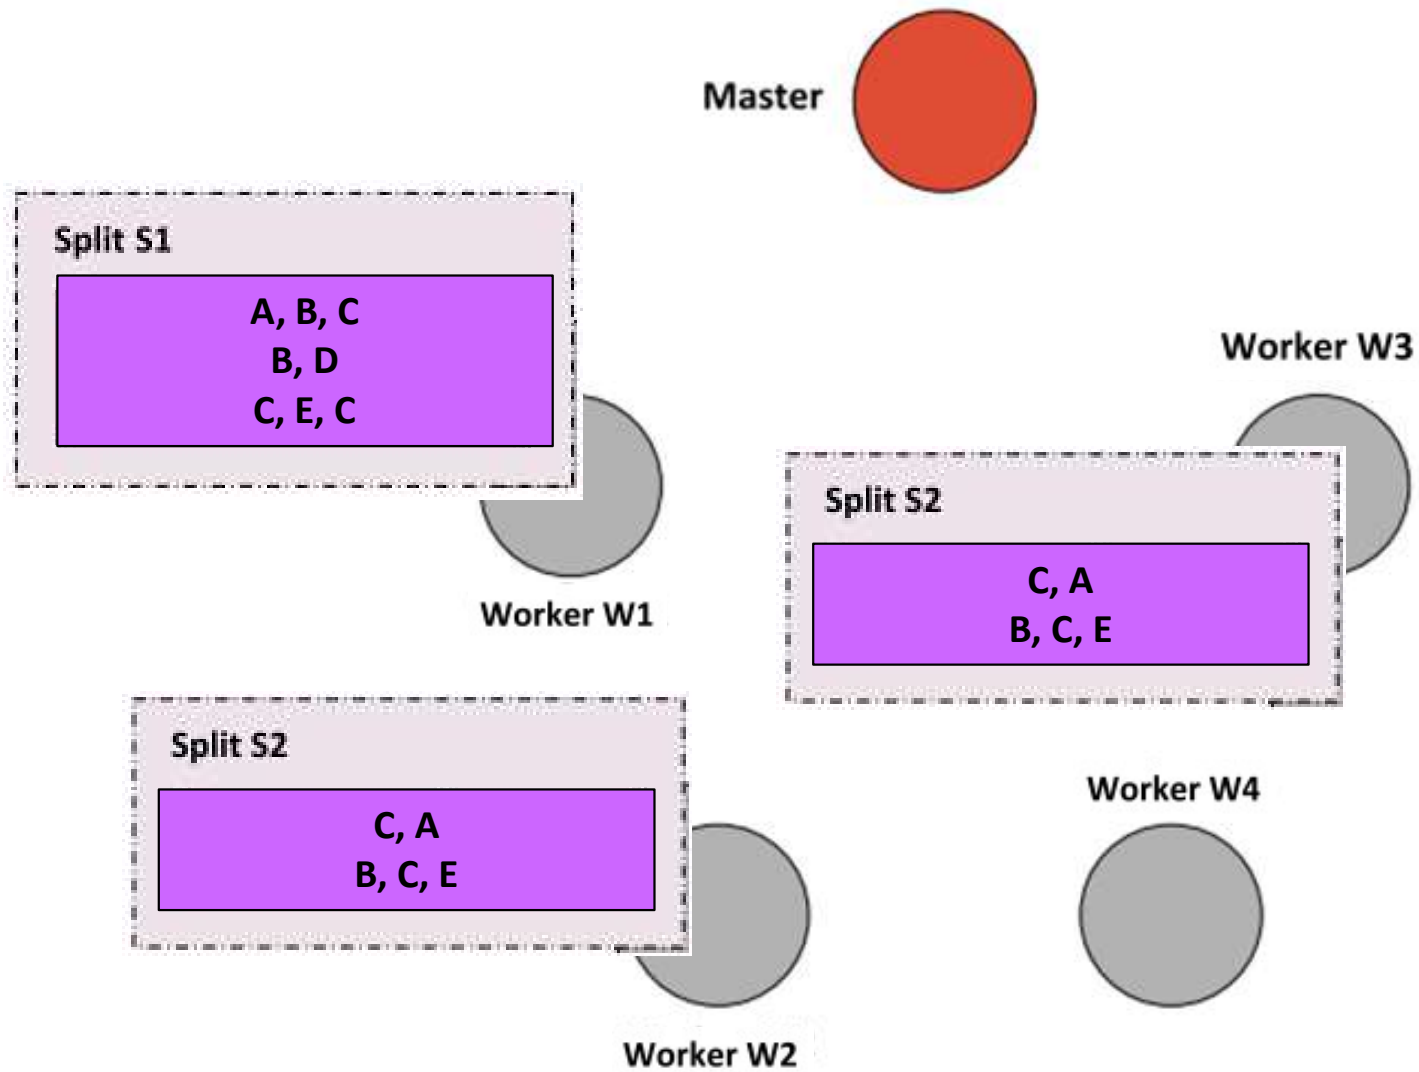
## Submission of MapReduce jobs

- Jobs can only be submitted to the master node
- Client provides the following:
  - **Implementation** of (not only) **Map and Reduce functions**
  - Description of **input file** (or even files)
  - Description of **output directory**

## Localization of input files

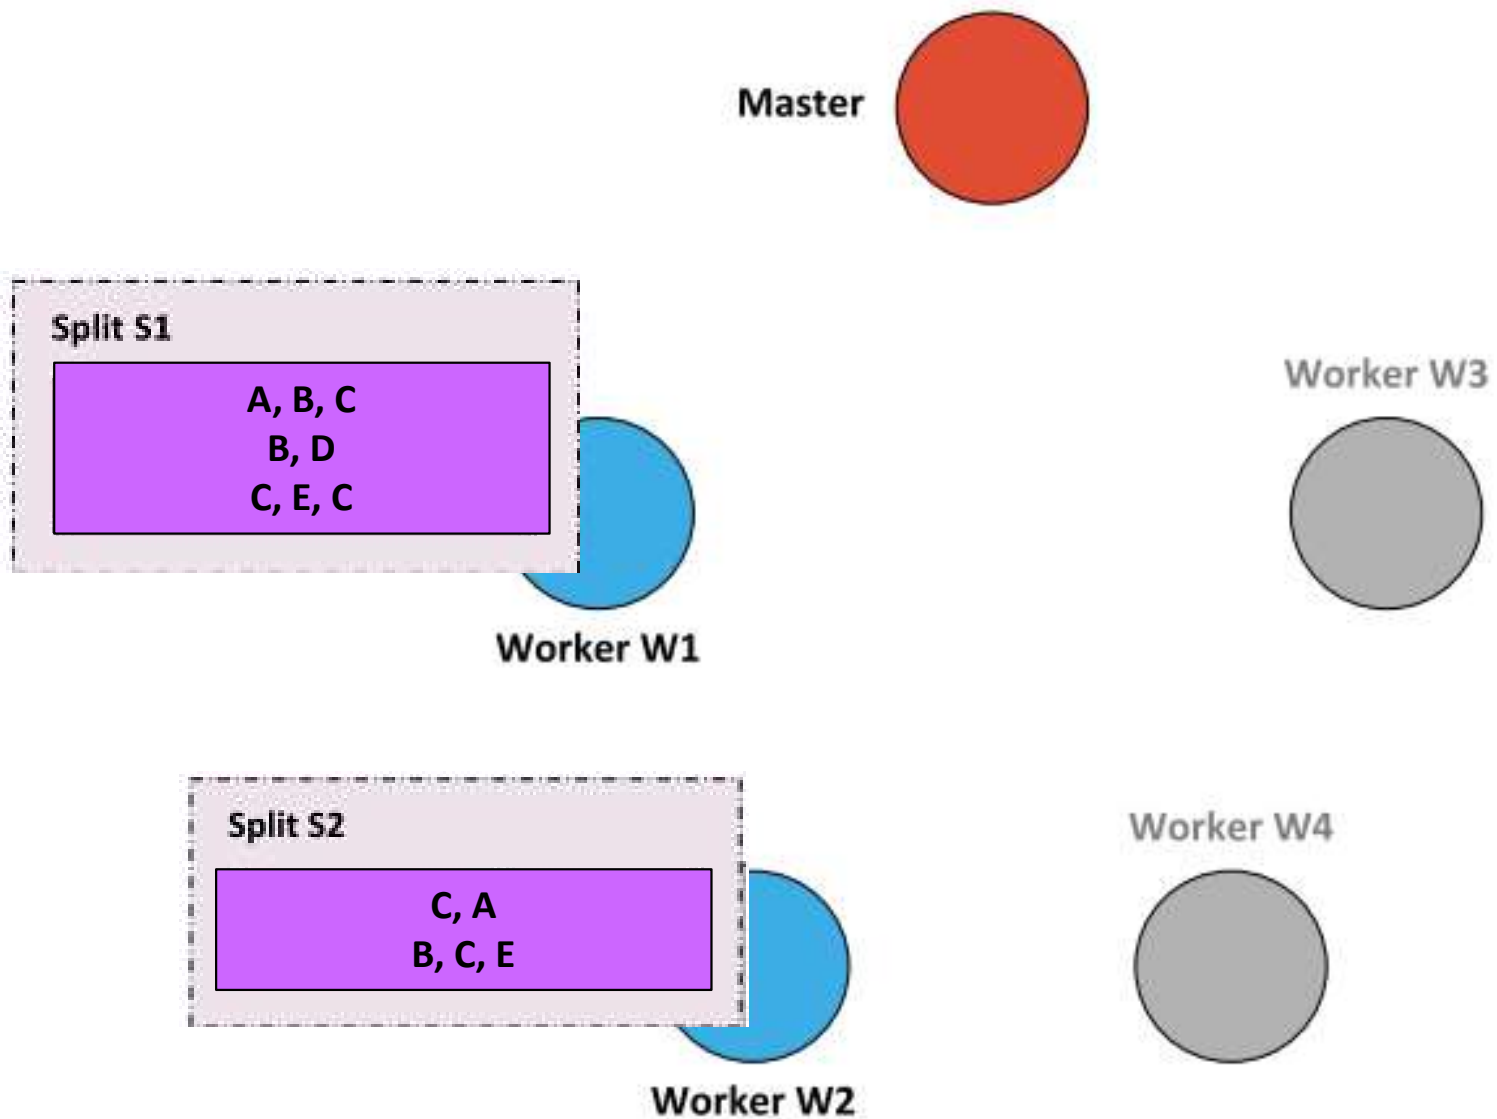
- Master determines **locations of all involved splits**
  - I.e. workers containing these splits are resolved

# Input Splits Localization

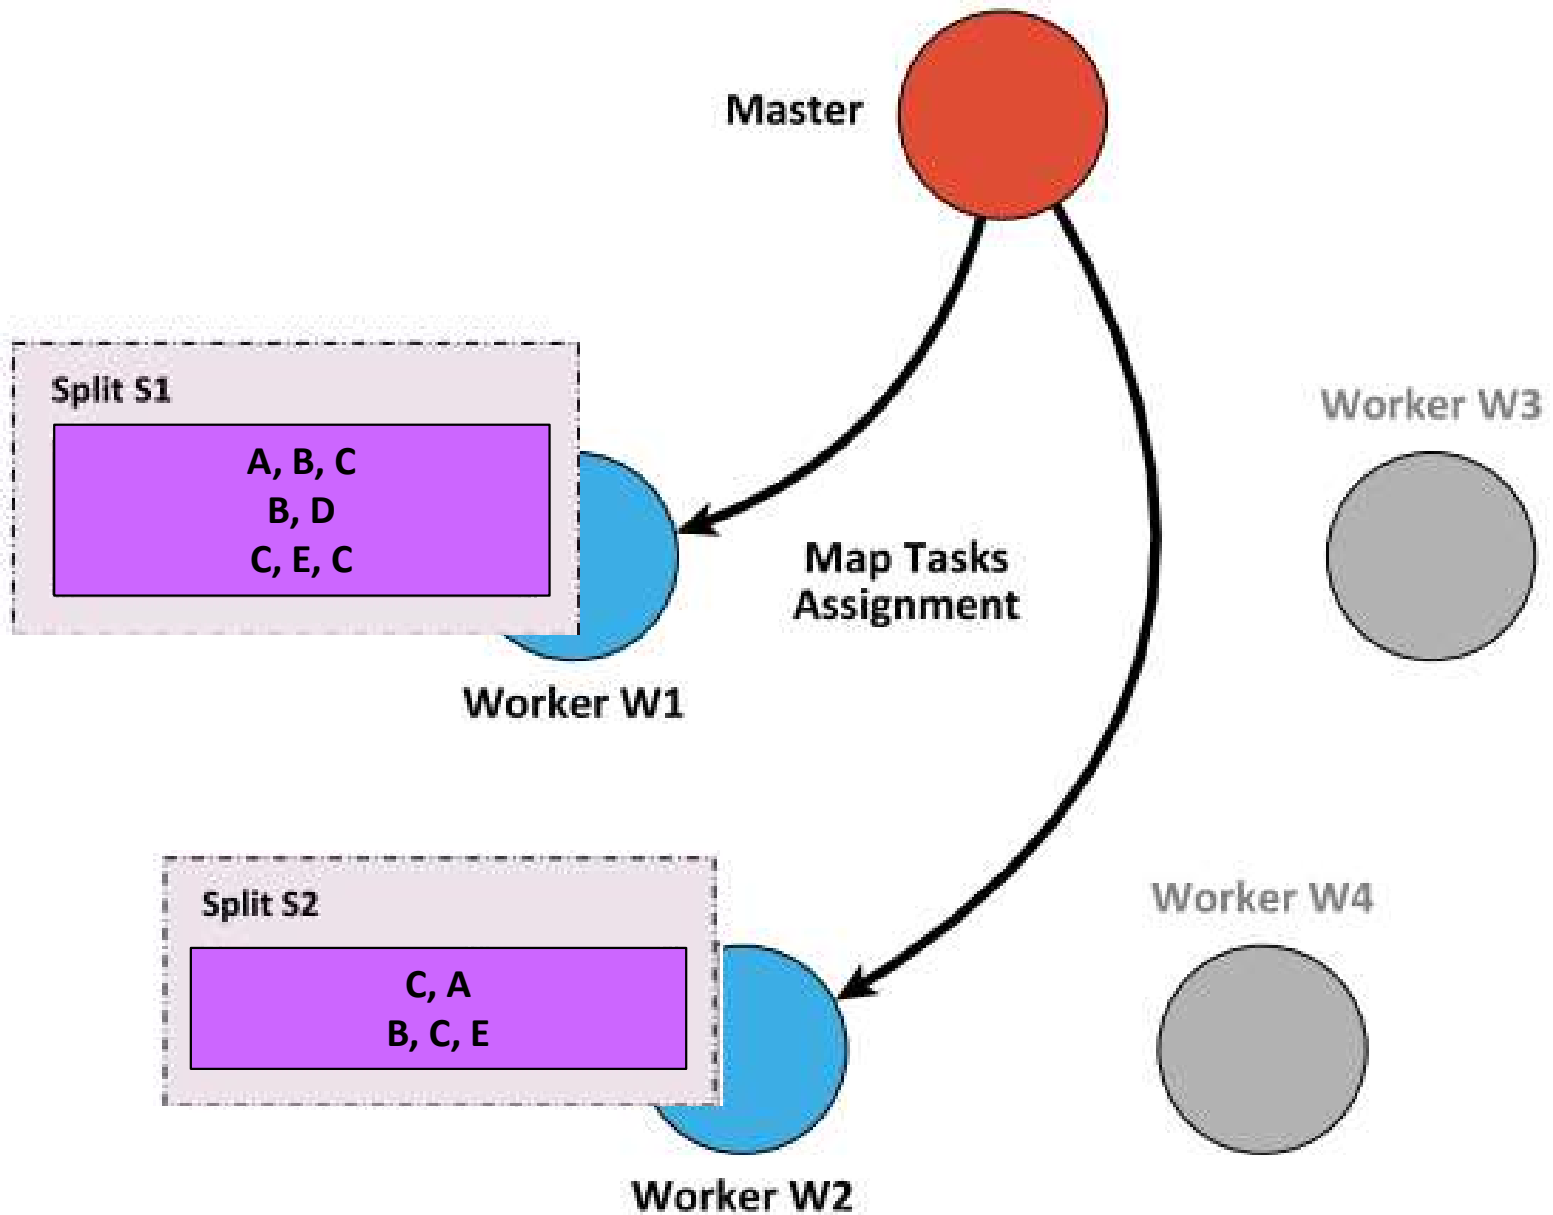




# Input Splits Localization



# Map Task Assignment



# Map Task Execution

**Map Task** = processing of 1 split by 1 worker

- Assigned by the master to an idle worker that is (preferably) already containing (physically storing) a given split

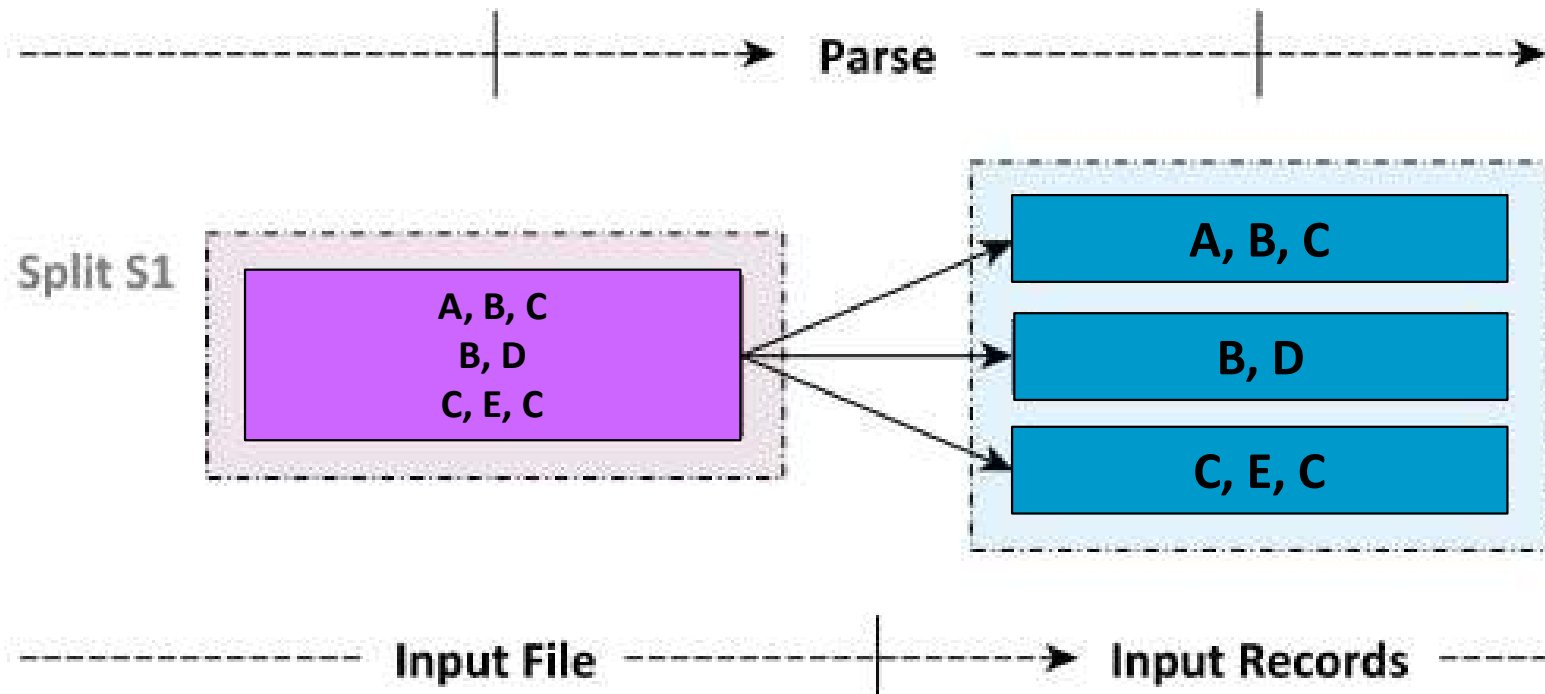
Individual steps...

- **Input reader** is used to **parse contents of the split**
  - I.e. **input records** are generated
- **Map function is applied on each input record**
  - Intermediate key-value pairs are emitted
- These pairs are **stored locally and organized into regions**
  - Either in the system memory,  
or flushed to a local hard drive when necessary
  - **Partition function** is used to determine the intended region
    - Intermediate keys (not values) are used for this purpose
    - E.g. hash of the key modulo the overall number of reducers

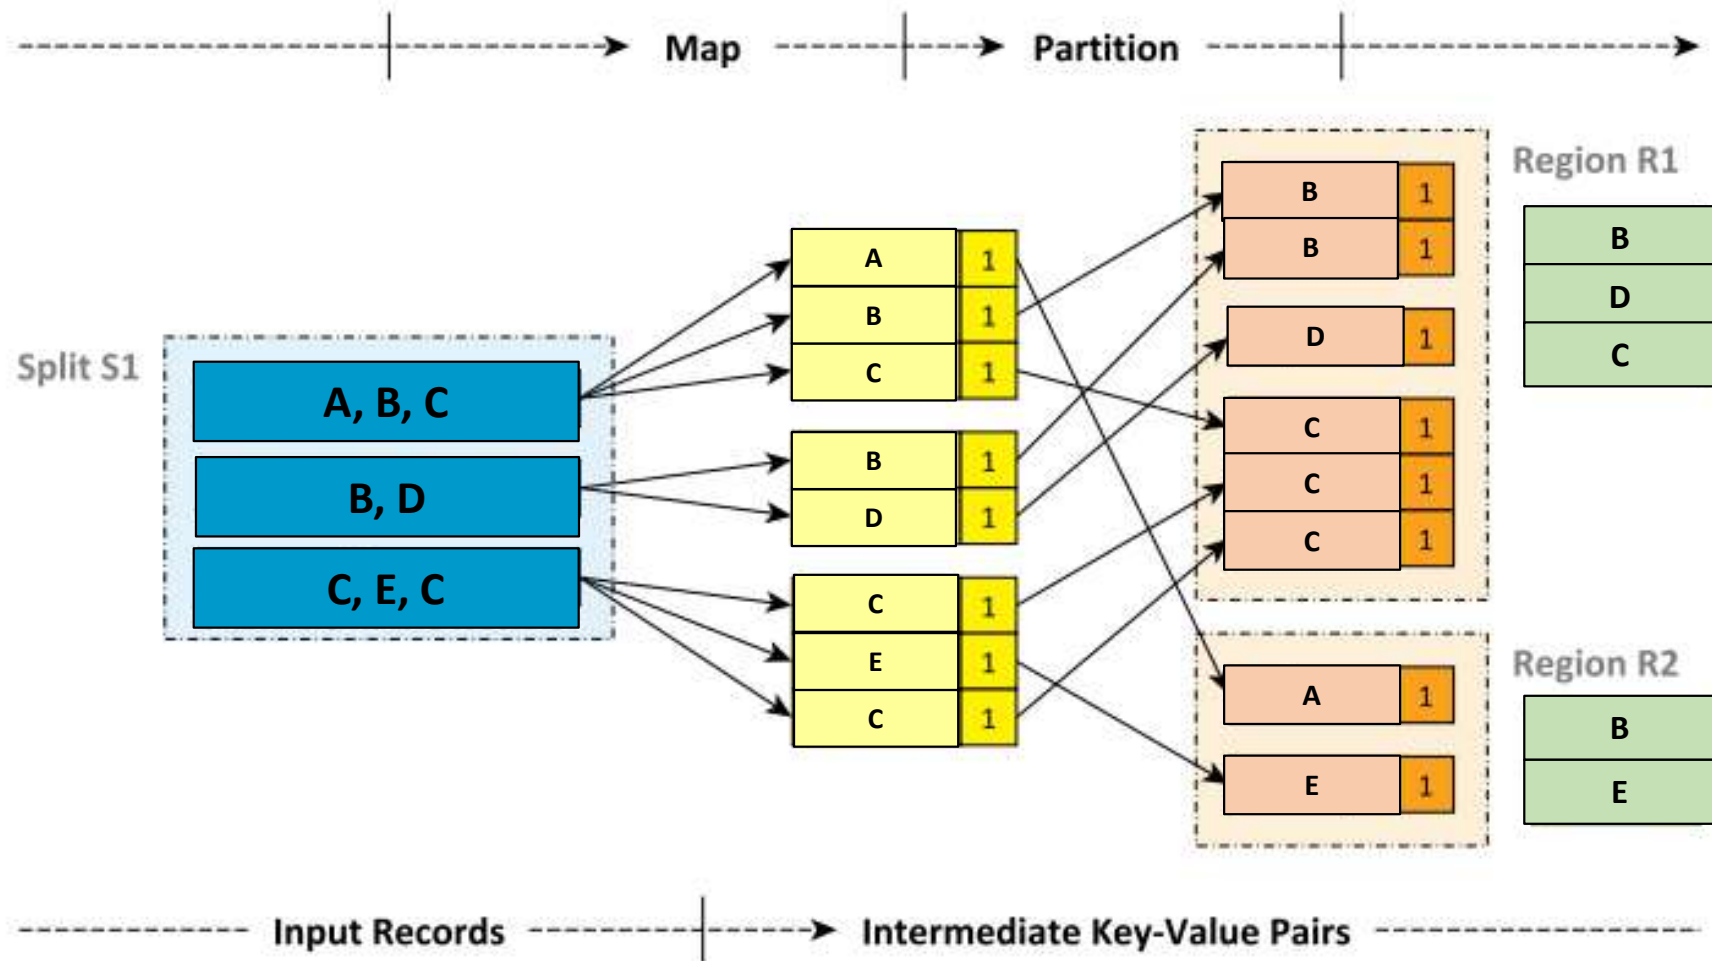
# Input Parsing

## Parsing phase

- **Each split is parsed** so that **input records** are retrieved (i.e. input key-value pairs are obtained)




# Map Phase






# Map Phase




Split S1





Master




Worker W3




Worker W1




Split S2




Worker W4




Worker W2



# Map Task Confirmation



# Reduce Task Assignment



# Reduce Task Execution

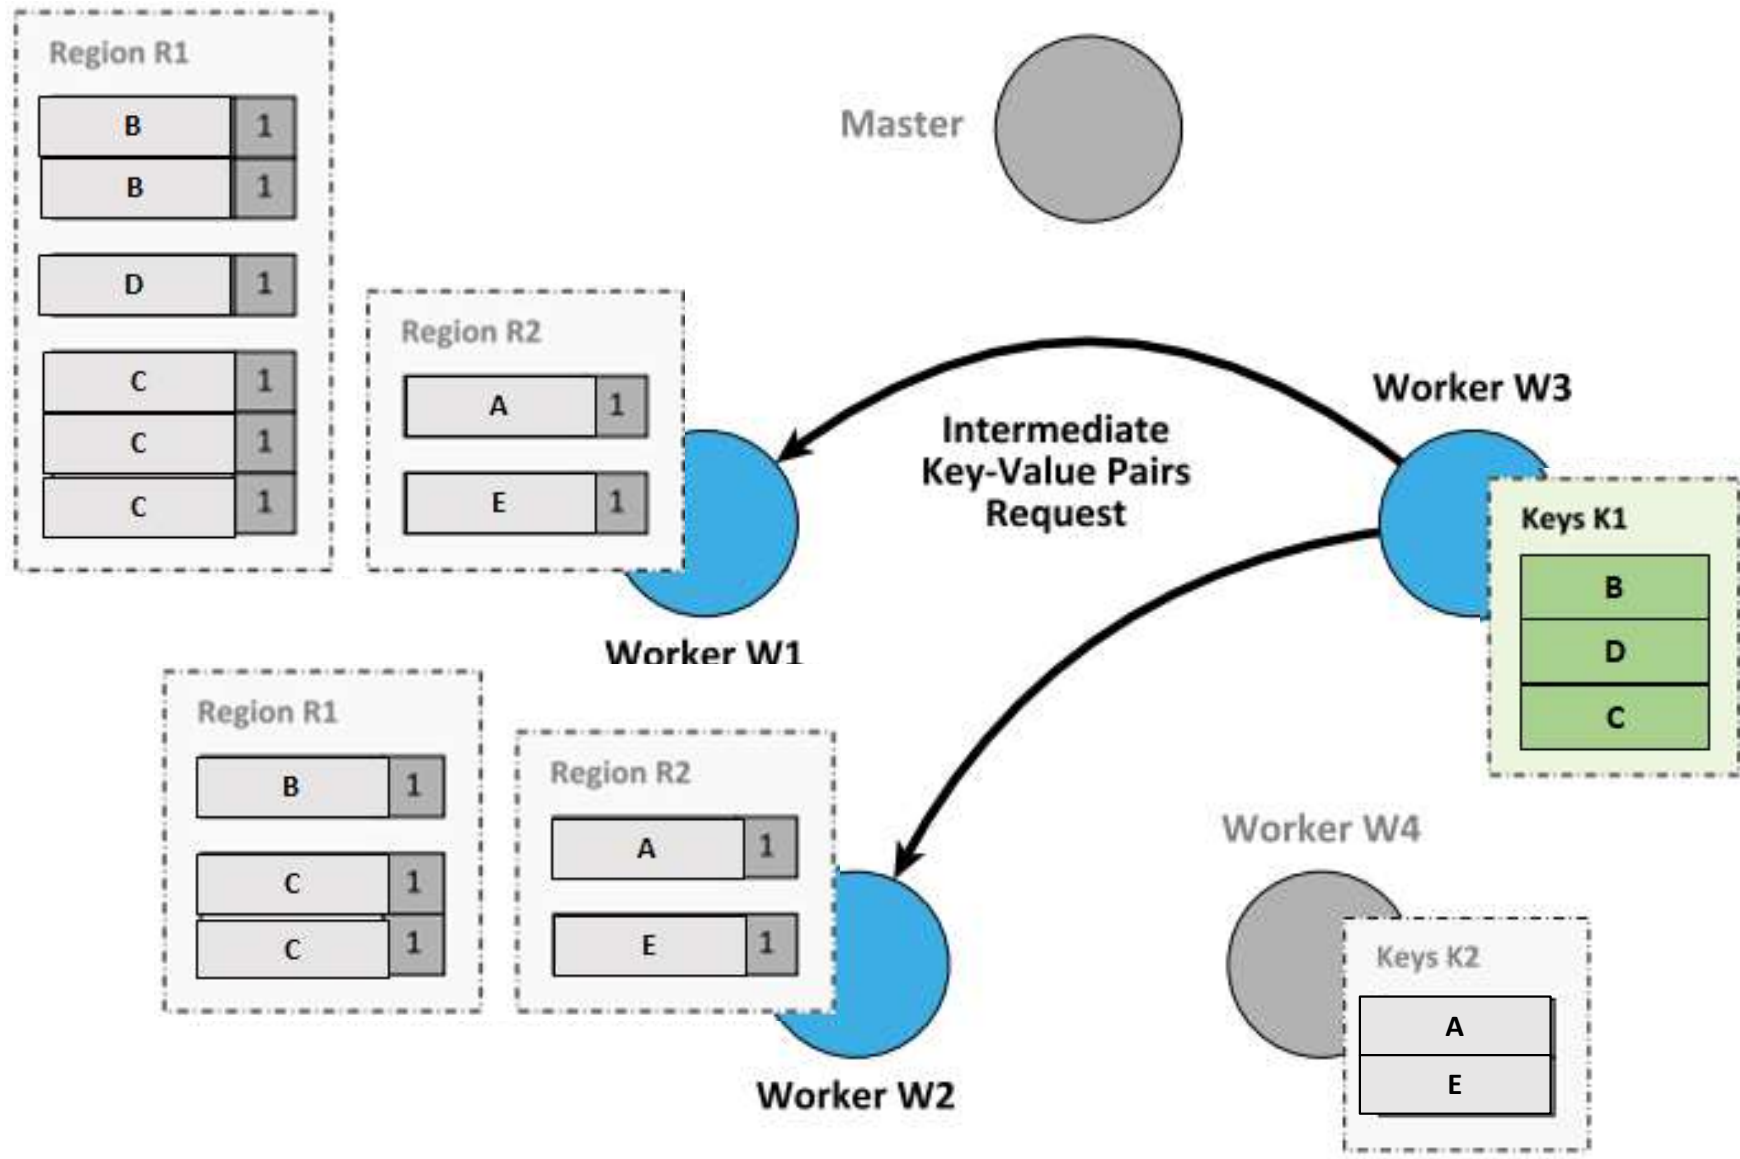
**Reduce Task** = reduction of selected key-value pairs by 1 worker

- Goal: processing of all emitted **intermediate key-value pairs belonging to a particular region**

Individual steps...

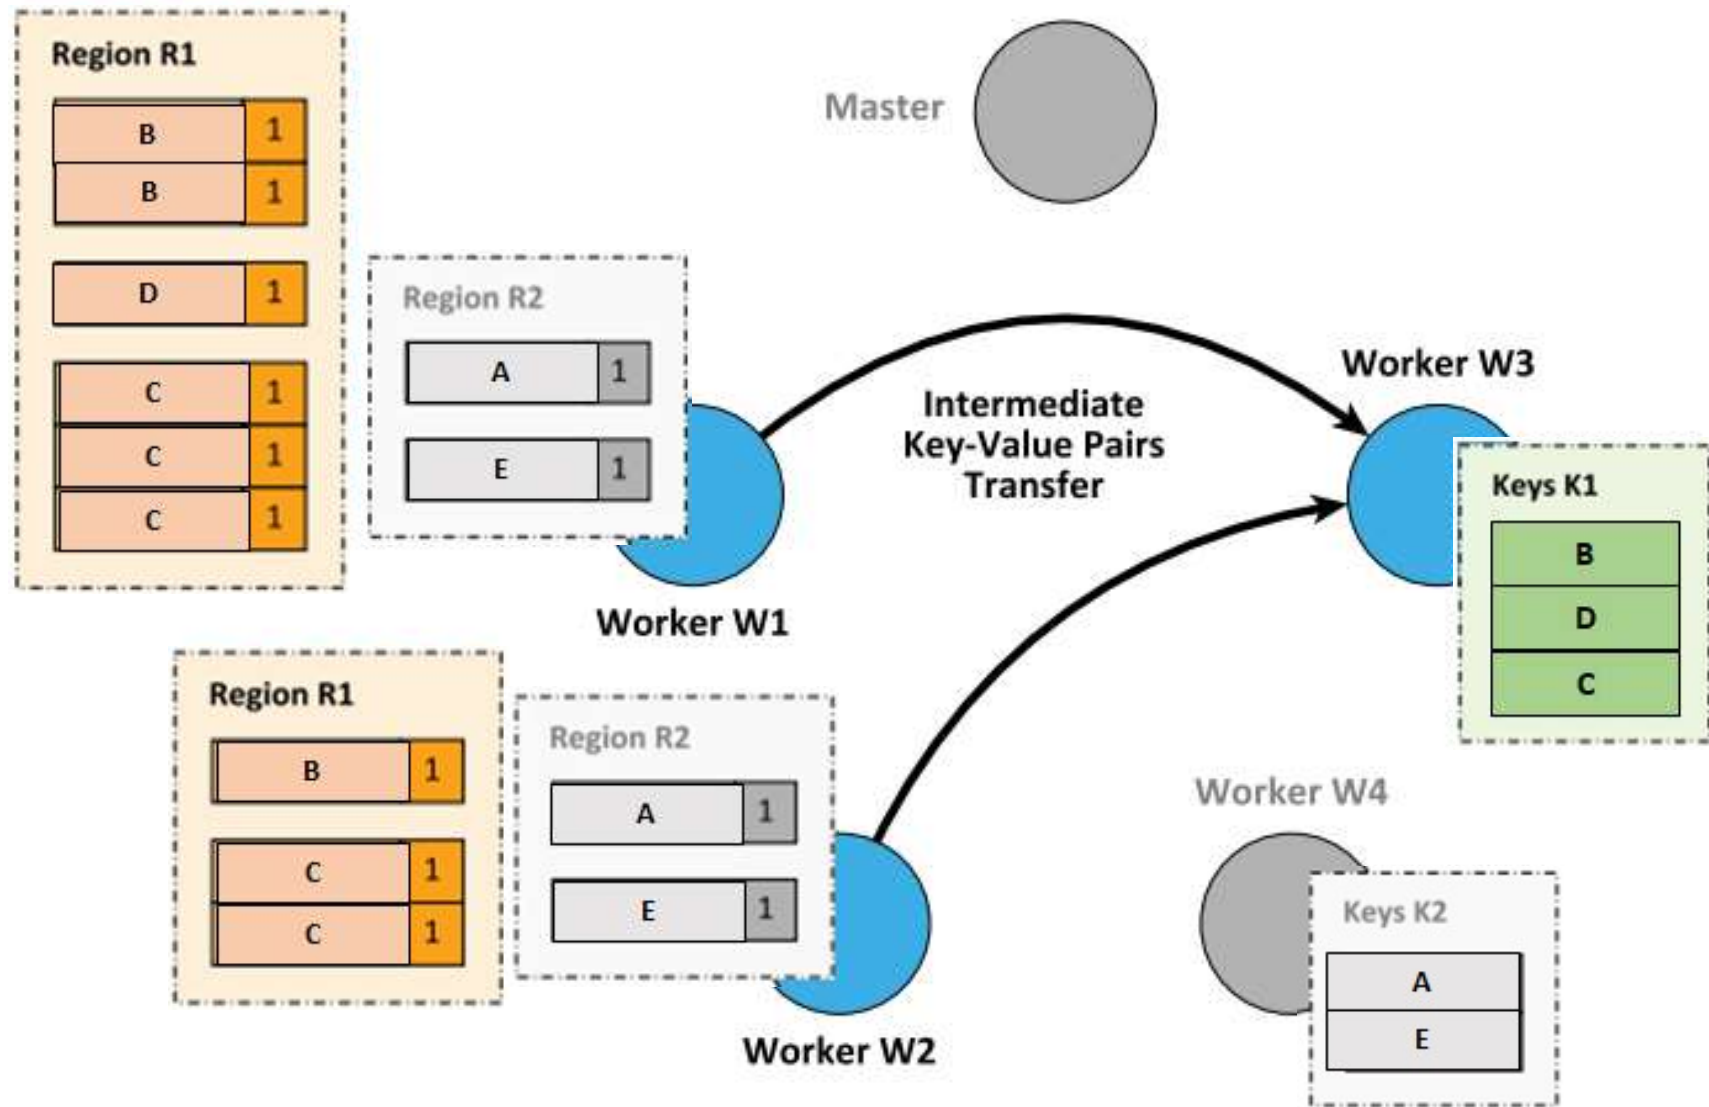
- **Intermediate key-value pairs are first acquired**
  - All relevant mapping workers are addressed
  - Data of corresponding **regions are transferred** (remote read)
- Once downloaded, they are **locally merged**
  - I.e. sorted and grouped based on keys
- **Reduce function** is applied on each intermediate key
- **Output key-value pairs** are emitted and stored (**output writer**)
  - Note that each worker produces its own separate output file

# Region Data Retrieval

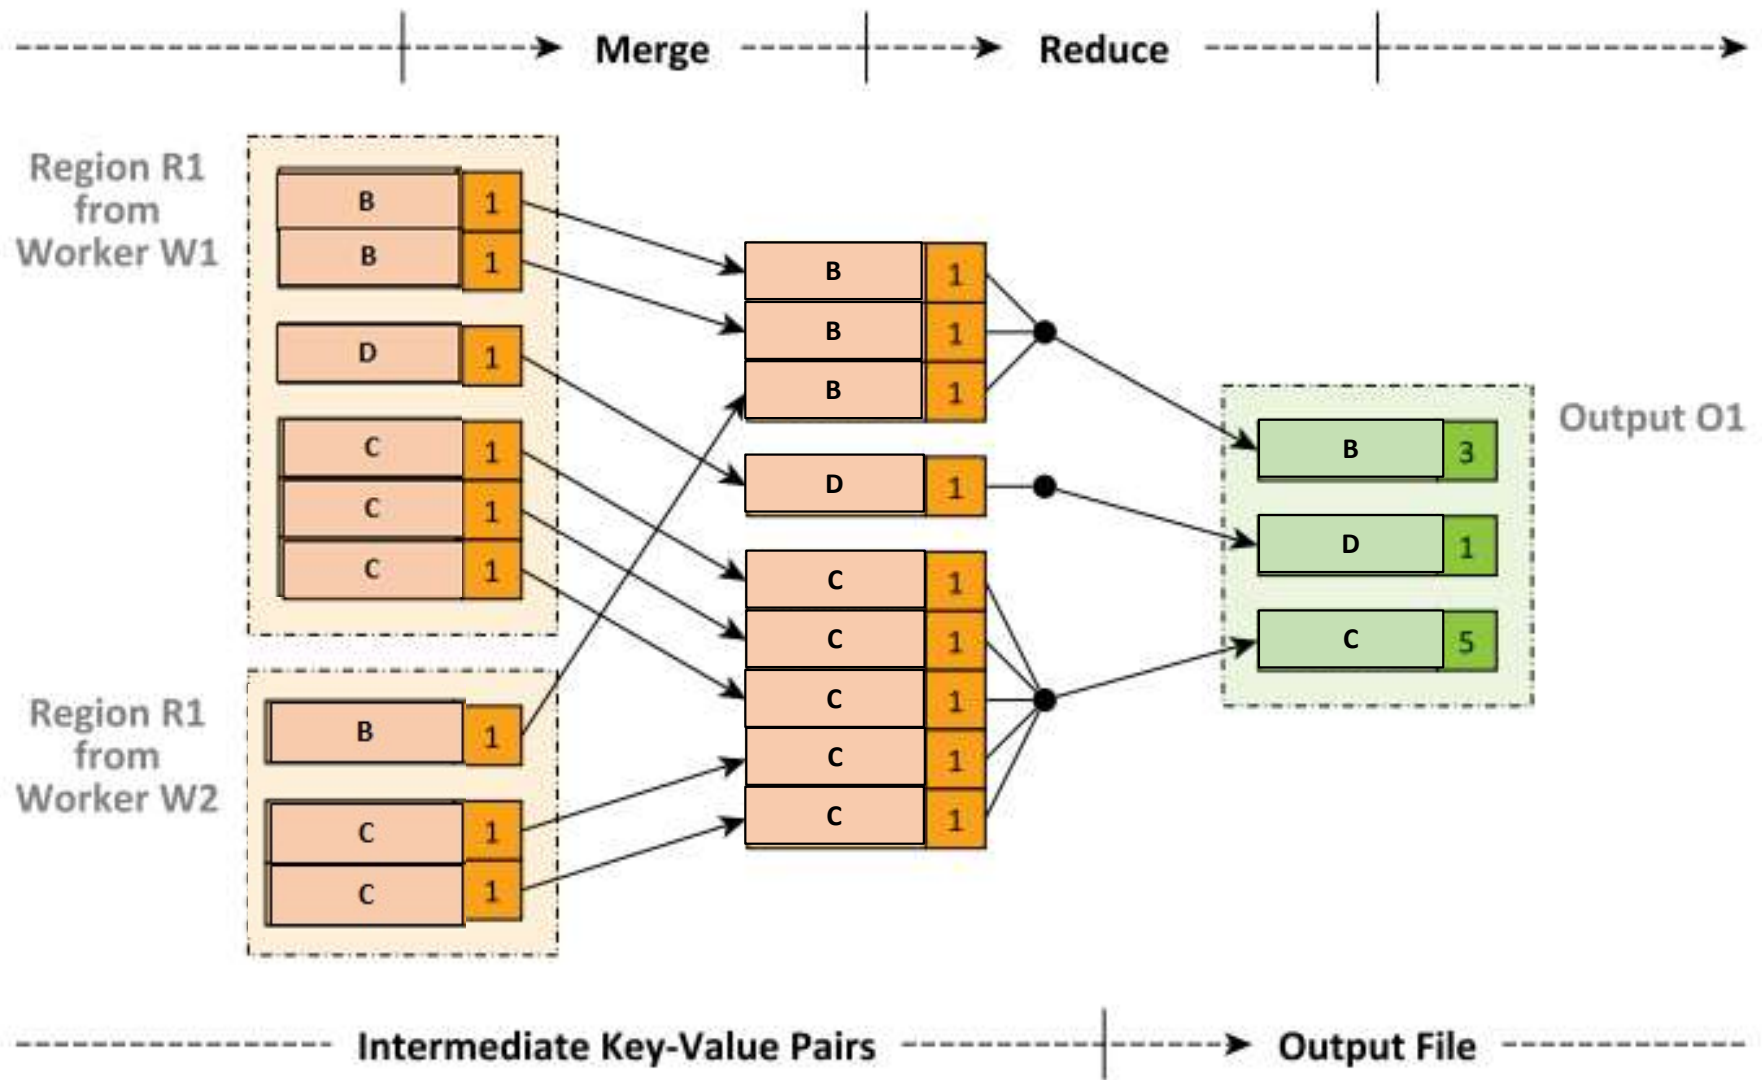




# Region Data Retrieval




# Reduce Phase



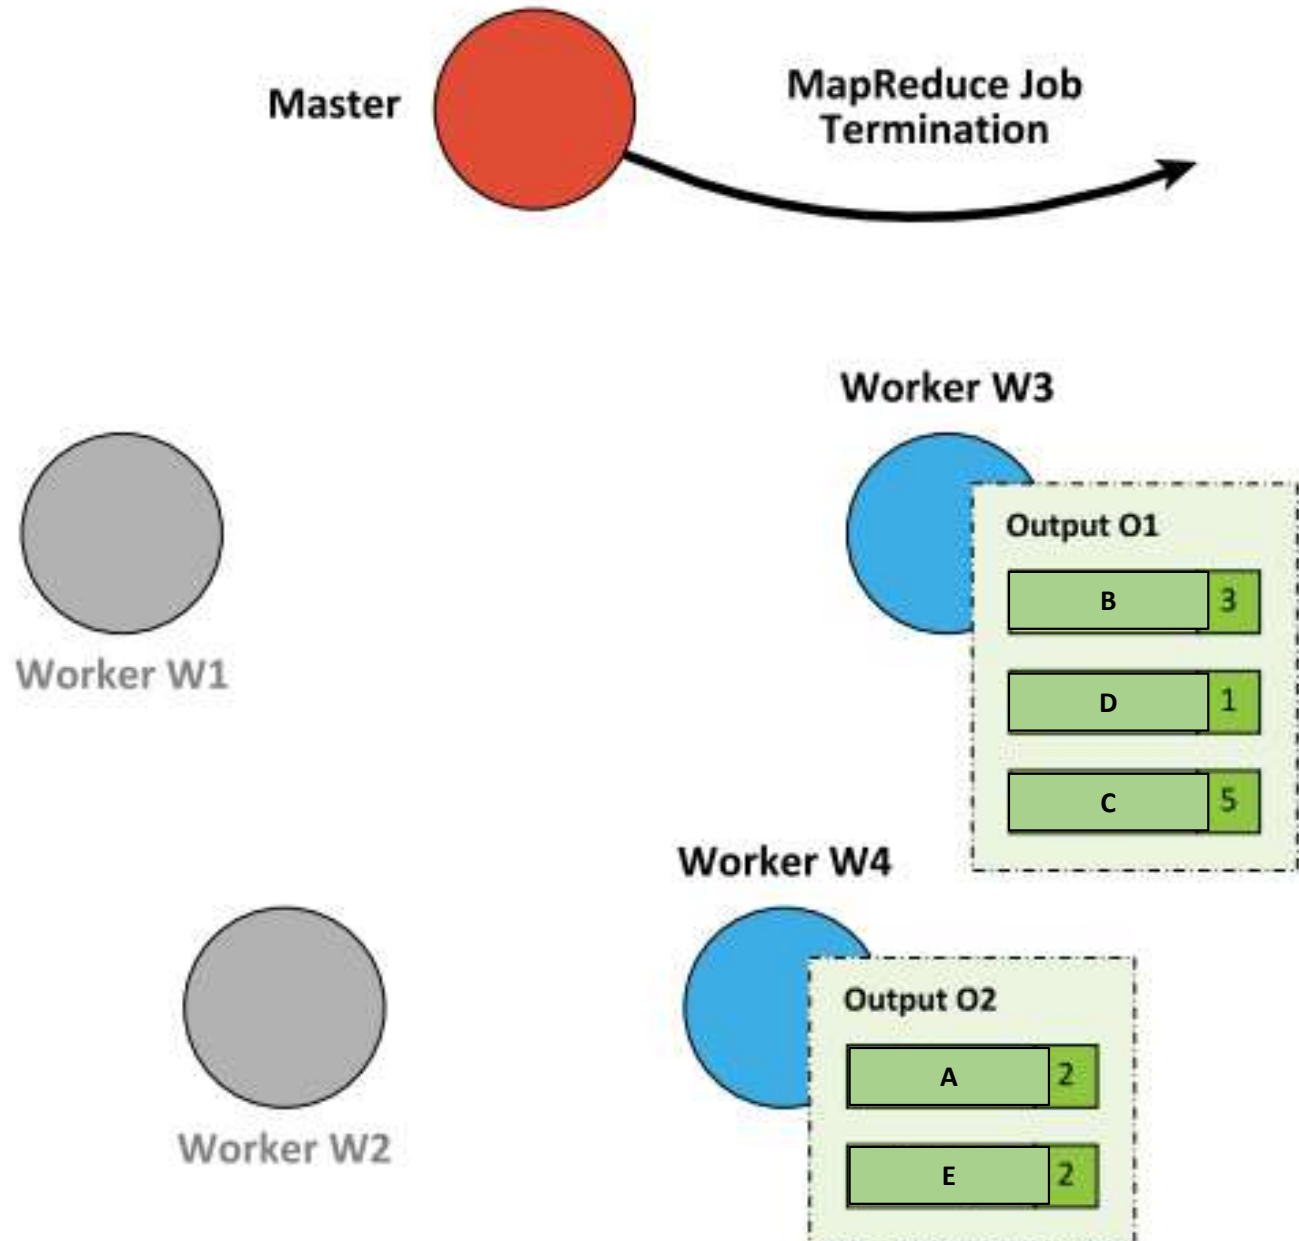
# Reduce Phase



# Reduce Task Confirmation



# MapReduce Job Termination



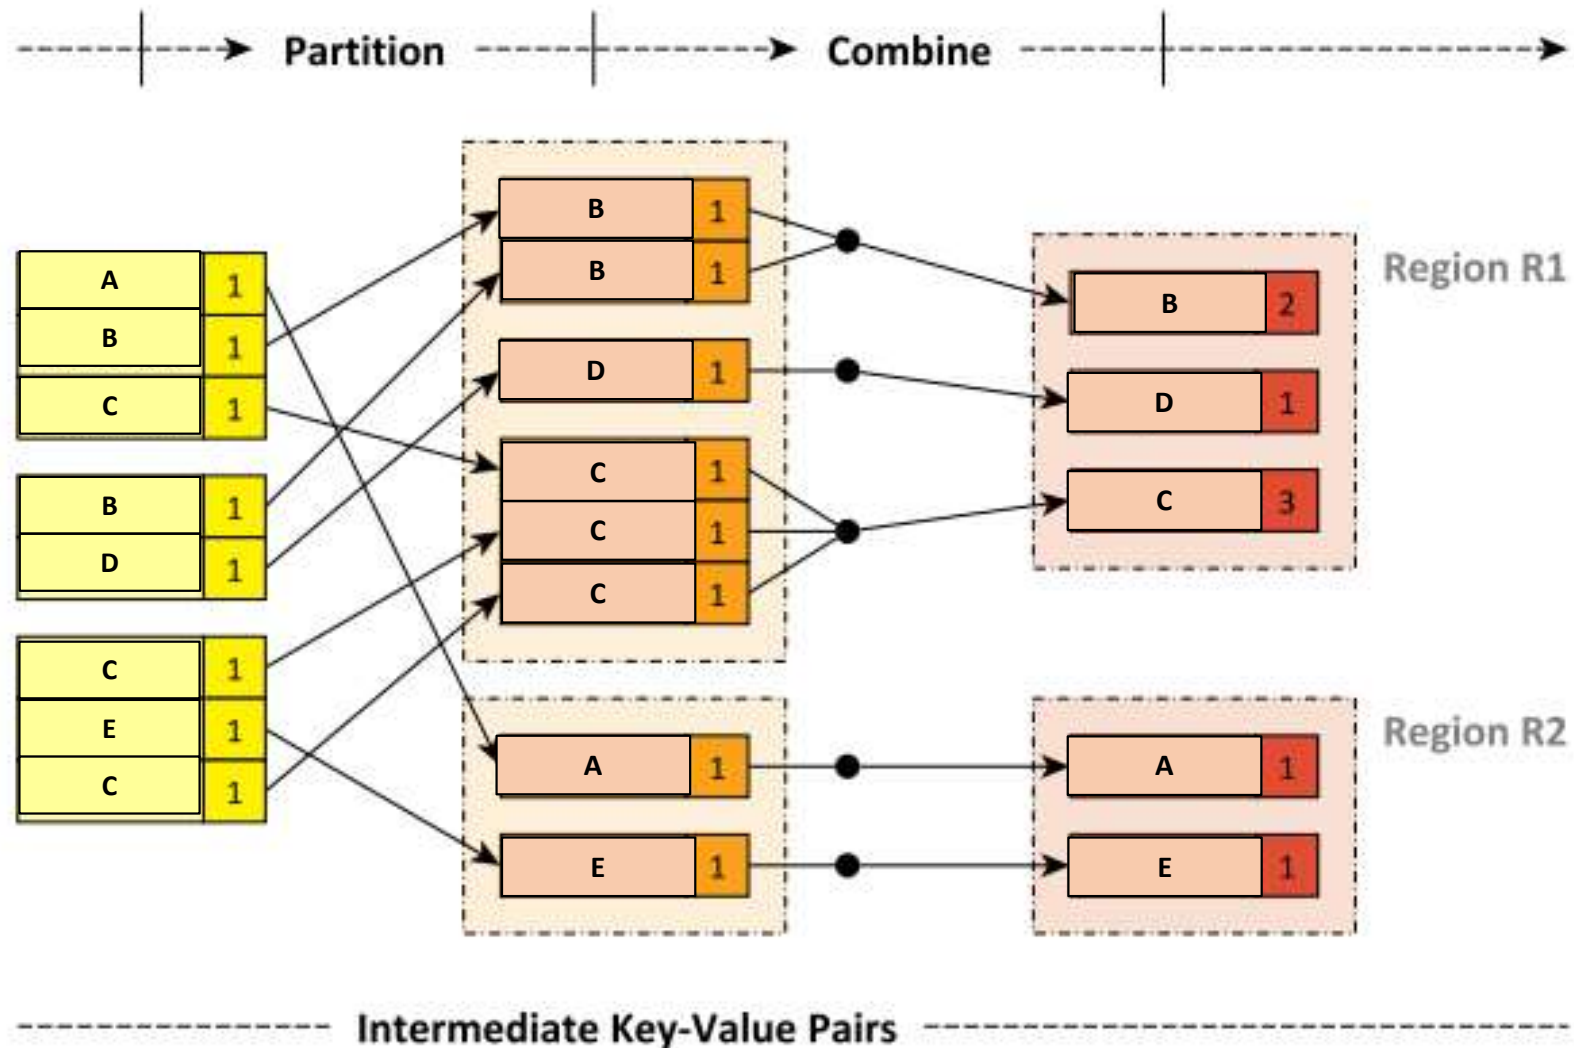
# Combine Function

## Optional **Combine** function


- Objective
  - **Decrease the amount of intermediate data**  
i.e. decrease the amount of data that is needed to be transferred from Mappers to Reducers
- Analogous purpose and implementation to **Reduce function**
- **Executed locally by Mappers**
- However, only applicable when the reduction is...
  - **Commutative**
  - **Associative**
  - **Idempotent:**  $f(f(x)) = f(x)$



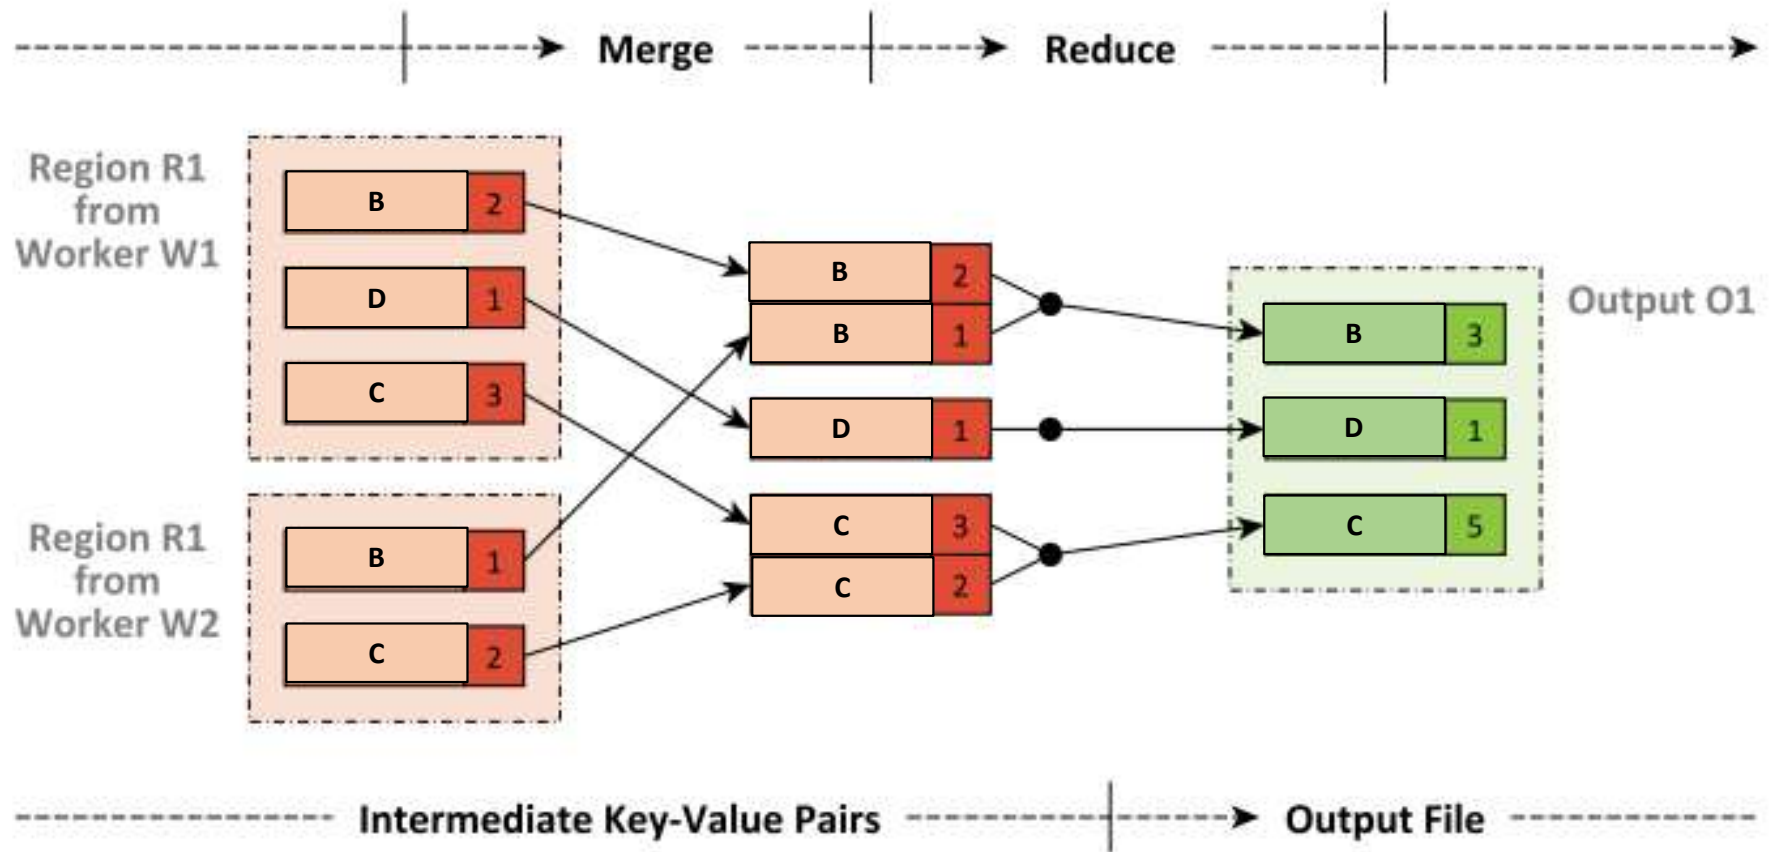
# Improved Map Phase



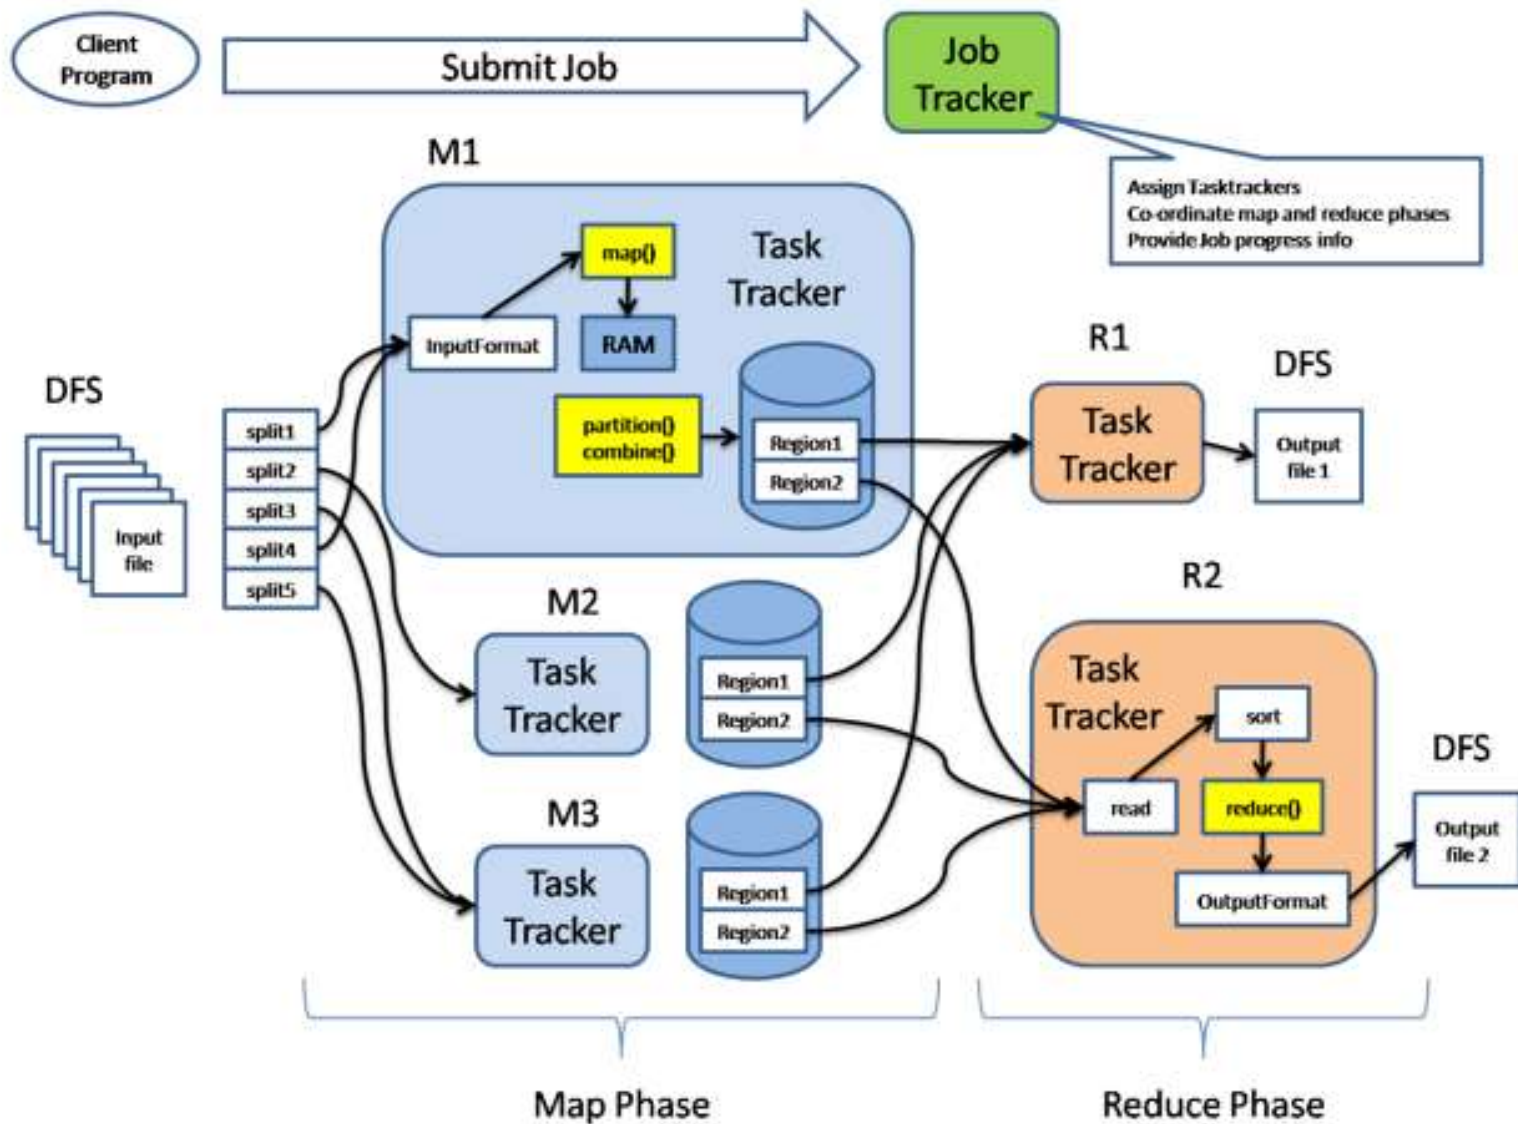
# Improved Map Phase



# Improved Map Phase



# Execution Schema



# Functions Overview

## Input reader

- Parses a given input split and **prepares input records**

## Map function

## Partition function

- **Determines a particular Reducer** for a given intermediate key

## Compare function

- Mutually **compares two intermediate keys**

## Combine function

## Reduce function

## Output writer

- **Writes the output** of a given Reducer



# Java Interface

## Mapper class

- Implementation of the **map function**
- Template parameters
  - KEYIN, VALUEIN – types of input key-value pairs
  - KEYOUT, VALUEOUT – types of intermediate key-value pairs
- Intermediate pairs are emitted via `context.write(k, v)`

```
class MyMapper extends Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {  
    @Override  
    public void map(KEYIN key, VALUEIN value, Context context)  
        throws IOException, InterruptedException  
    {  
        // Implementation  
    }  
}
```



# Java Interface

## Reducer class

- Implementation of the **reduce function**
- Template parameters
  - KEYIN, VALUEIN – types of intermediate key-value pairs
  - KEYOUT, VALUEOUT – types of output key-value pairs
- Output pairs are emitted via `context.write(k, v)`

```
class MyReducer extends Reducer<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {  
    @Override  
    public void reduce(KEYIN key, Iterable<VALUEIN> values, Context context)  
        throws IOException, InterruptedException  
    {  
        // Implementation  
    }  
}
```

# Example

## Word Frequency

- *Input*: Documents with words
  - Files located at `/home/input` HDFS directory
- *Map*: parses a document, emits (word, 1) pairs
- *Reduce*: computes and emits the sum of the associated values
- *Output*: overall number of occurrences for each word
  - Output will be written to `/home/output`

## MapReduce **job execution**

```
hadoop jar wc.jar WordCount /home/input /home/output
```

# Example: Mapper Class

```
public class WordCount {  
    ...  
    public static class MyMapper  
        extends Mapper<Object, Text, Text, IntWritable>  
    {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
        @Override  
        public void map(Object key, Text value, Context context)  
            throws IOException, InterruptedException  
        {  
            StringTokenizer itr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
    ...  
}
```

# Example: Reducer Class

```
public class WordCount {  
    ...  
    public static class MyReducer  
        extends Reducer<Text, IntWritable, Text, IntWritable>  
    {  
        private IntWritable result = new IntWritable();  
        @Override  
        public void reduce(Text key, Iterable<IntWritable> values,  
            Context context) throws IOException, InterruptedException  
        {  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            result.set(sum);  
            context.write(key, result);  
        }  
    }  
    ...  
}
```

# Advanced Aspects

## Counters

- Allow to track the progress of a MapReduce job in real time
  - **Predefined counters**
    - E.g. numbers of launched / finished Map / Reduce tasks, parsed input key-value pairs, ...
  - **Custom counters** (user-defined)
    - Can be associated with any action that a Map or Reduce function does

## Fault tolerance

- When a large number of nodes process a large number of data  
⇒ **fault tolerance is necessary**



# Advanced Aspects

## Worker failure

- Master periodically pings every worker; if no response is received in a certain amount of time, master marks the worker as failed
- **All its tasks are reset back to their initial idle state and become eligible for rescheduling on other workers**

## Master failure

- Strategy A – periodic checkpoints are created; if master fails, a new copy can then be started
- Strategy B – master failure is considered to be highly unlikely; users simply resubmit unsuccessful jobs



# Advanced Aspects

## Stragglers

- **Straggler** = node that takes unusually long time to complete a task it was assigned
- Solution
  - When a MapReduce job is close to completion, the master schedules **backup executions** of the remaining in-progress tasks
  - A given task is considered to be completed whenever either the primary or the backup execution completes

# Additional Examples

## URL access frequency

- *Input*: HTTP server access logs
- *Map*: parses a log, emits (accessed URL, 1) pairs
- *Reduce*: computes and emits the sum of the associated values
- *Output*: overall number of accesses to a given URL

## Inverted index

- *Input*: text documents containing words
- *Map*: parses a document, emits (word, document ID) pairs
- *Reduce*: emits all the associated document IDs sorted
- *Output*: list of documents containing a given word

# Additional Examples

## Distributed sort

- *Input*: records to be sorted according to a specific criterion
- *Map*: extracts the sorting key, emits (key, record) pairs
- *Reduce*: emits the associated records unchanged

## Reverse web-link graph

- *Input*: web pages with `<a href="...">...</a>` tags
- *Map*: emits (target URL, current document URL) pairs
- *Reduce*: emits the associated source URLs unchanged
- *Output*: list of URLs of web pages targeting a given one

# Additional Examples

## Reverse web-link graph

```
/**
 * Map function
 * @param key    Source web page URL
 * @param value  HTML contents of this web page
 */
map(String key, String value) {
    foreach <a> tag t in value: emit(t.href, key);
}
```

```
/**
 * Reduce function
 * @param key    URL of a particular web page
 * @param values List of URLs of web pages targeting this one
 */
reduce(String key, Iterator values) {
    emit(key, values);
}
```

# Use Cases: General Patterns

## **Counting, summing, aggregation**

- When the overall number of occurrences of certain items or a different aggregate function should be calculated

## **Collating, grouping**

- When all items belonging to a certain group should be found, collected together or processed in another way

## **Filtering, querying, parsing, validation**

- When all items satisfying a certain condition should be found, transformed or processed in another way

## **Sorting**

- When items should be processed in a particular order with respect to a certain ordering criterion

# Use Cases: Real-World Problems

Just a few **real-world examples...**

- Risk modeling, customer churn
- Recommendation engine, customer preferences
- Advertisement targeting, trade surveillance
- Fraudulent activity threats, security breaches detection
- Hardware or sensor network failure prediction
- Search quality analysis
- ...



# Real world Problems

Use Case	Description
<b>1. Word Count / Log Analysis</b>	Counting frequency of words or log entries from large-scale documents/logs.
<b>2. Indexing Web Pages</b>	Creating inverted indexes (like in search engines) from web crawled data.
<b>3. Recommendation Engines</b>	Used in collaborative filtering for e-commerce/movie platforms.
<b>4. Sentiment Analysis</b>	Extracting and aggregating sentiments from large text datasets (e.g., tweets, reviews).
<b>5. Network Traffic Monitoring</b>	Analyzing logs from distributed network devices for trends or threats.
<b>6. Genomics / DNA Processing</b>	Processing and comparing millions of DNA sequences efficiently.
<b>7. Financial Risk Modeling</b>	Large-scale simulation and aggregation of financial datasets.

## MapReduce criticism

- MapReduce **is a step backwards**
  - Does not use database schema
  - Does not use index structures
  - Does not support advanced query languages
  - Does not support transactions, integrity constraints, views, ...
  - Does not support data mining, business intelligence, ...
- MapReduce **is not novel**
  - Ideas more than 20 years old and overcome
  - Message Passing Interface (MPI), Reduce-Scatter

The end of MapReduce?