



Natural Language Processing (NLP)

N Grams & Intro of Embedding

Equipping You with Research Depth and
Industry Skills – Data Science Oriented

By:

Dr. Zohair Ahmed



www.youtube.com/@ZohairAI

Subscribe



www.begindiscovery.com

Introduction to Text Representation

- **Text Representation:** The process of converting text into numerical form for machine learning tasks.
- **Importance in NLP:** Text must be represented as vectors to be processed by algorithms.
- **Common Methods:**
 - Bag of Words (BoW)
 - TF-IDF
 - Word Embeddings
 - n-grams



What is an n-gram?

- **Definition:** A sequence of n consecutive words in a text.
- **Example:**
 - Unigrams (1-gram): "I", "am", "learning"
 - Bigrams (2-grams): "I am", "am learning"
 - Trigrams (3-grams): "I am learning"
- **Why n-grams?:** Captures context and word relationships beyond single words.



Bag of n-grams Overview

- **Bag of n-grams:** An extension of the Bag of Words (BoW) model.
 - Represents a text by counting occurrences of n-grams (not individual words).
- **Features:**
 - Can capture context between consecutive words.
 - Works well for tasks where word order matters (e.g., text classification, language modeling).



How Bag of n-grams Works

- **Step 1:** Split the text into sentences and then into n-grams.
 - Example: "I am learning NLP."
 - Bigrams: ("I am", "am learning", "learning NLP")
- **Step 2:** Count the frequency of each n-gram across the entire corpus.
- **Step 3:** Represent the text as a vector of n-gram frequencies.



Example of Bag of n-grams

- **Text:** "I love programming and I love learning."
- **Unigrams:** {"I", "love", "programming", "and", "learning"}
- **Bigrams:** {"I love", "love programming", "programming and", "and I", "I love", "love learning"}
- **Trigrams:** {"I love programming", "love programming and", "programming and I", "and I love", "I love learning"}

Advantages of Bag of n-grams

- **Captures Context:** Unlike BoW, it can model relationships between consecutive words.
- **Improved Performance:** Especially useful for tasks like language modeling, sentiment analysis, and text classification.
- **Simplicity:** Easy to implement and understand.



Challenges of Bag of n-grams

- **Increased Dimensionality:** The number of possible n-grams can grow exponentially with large corpora.
- **Sparsity:** Many n-grams may not appear frequently enough to be useful.
- **Ignoring Word Order:** Although n-grams capture some context, they still do not fully capture complex word relationships beyond n-gram length.

Applications of Bag of n-grams

- **Text Classification:** Categorizing text into predefined categories.
- **Sentiment Analysis:** Identifying the sentiment (positive/negative) of a piece of text.
- **Machine Translation:** Translating sentences based on learned n-grams.
- **Speech Recognition:** Modeling language patterns in spoken language.



Applications of Bag of n-grams

- **Text Classification:** Categorizing text into predefined categories.
- **Sentiment Analysis:** Identifying the sentiment (positive/negative) of a piece of text.
- **Machine Translation:** Translating sentences based on learned n-grams.
- **Speech Recognition:** Modeling language patterns in spoken language.



Word Embeddings



Limitations of Bag of Words & TF-IDF

- **Large Vector Sizes:**
 - Example: Vocabulary of 200,000 words, resulting in 100,000-size vectors.
 - **High Memory & Compute Usage.**
- **Sparsity:**
 - Most values are zeros, leading to inefficient representations.
- **Example:**
 - Sentences like "I need help" vs. "I need assistance" might not have similar vector representations in TF-IDF or Bag of Words, despite being semantically similar.

What is Word Embedding?

- **Word Embedding:** A technique to represent words in dense vectors with a lower dimensionality.
- **Key Advantage:**
 - Similar words (like "good" and "great") have **similar vectors**.
 - Word embeddings have **dense representations** (fewer zeros).
- **Vector Size:** Typically around **300 dimensions**, which is much smaller than traditional models.

How Word Embedding Works

- **Similar Words = Similar Vectors:**
 - Example: "Good" and "Great" will have vectors like:
 - a. "Good": [3.1, 3.1, 4.4, 4.2]
 - b. "Great": [3.2, 3.1, 4.3, 4.1]
- **Efficient:** Smaller vectors, fewer zeros, and captures the meaning of words.



Word Embedding Techniques

- **Popular Techniques:**
 - Word2Vec
 - GloVe
 - FastText
- **Word2Vec:**
 - Uses **Continuous Bag of Words (CBOW)** and **Skip-gram**.
- **Transformer-based Models:**
 - **BERT, GPT:** Recent advancements that improve text representation further.



The Power of Word2Vec

- **Word2Vec Example:**
 - **Arithmetic with Words:**
 - a. King - Man + Woman = Queen
 - b. Word2Vec can perform **arithmetic** with word vectors to capture relationships.
- **Visualizing Word Relationships:**
 - **King** and **Queen** share common attributes like authority and power, but differ by gender.



Dataset Variations in Word Embeddings

- **Word2Vec**: Train on different corpora to get domain-specific embeddings.
 - Example: Google News, Amazon Reviews, etc.
- **GloVe**: Can be trained on specific datasets like Twitter or Wikipedia.
 - Twitter dataset understands slang and abbreviations better.
- **BERT**: Can be fine-tuned on specific domains (e.g., BioBERT for biomedical data, FinBERT for financial data).



Other Advanced Techniques

- **Elmo:** Embeddings based on LSTM (Long Short-Term Memory).
- **Transformer Models:**
 - BERT, Albert, Roberta, etc.
 - Trained on specific datasets for better domain understanding.



Converting Text to Vectors

- **Goal:** Convert individual words, sentences, or entire documents into vectors.
 - **Sentence Embedding:** Converting a sentence into a single vector.
 - **Document Embedding:** Converting a document or article into a vector.
- **Why:** Machine learning models need numbers, and word embeddings allow text to be represented in a numerical format that captures meaning.

