



# Natural Language Processing (NLP)

## Sequence to Sequence Architecture

Equipping You with Research Depth and  
Industry Skills

By:

Dr. Zohair Ahmed



[www.youtube.com/@ZohairAI](https://www.youtube.com/@ZohairAI)

**Subscribe**



[www.begindiscovery.com](http://www.begindiscovery.com)

# Sequence-to-Sequence (Seq2Seq)

---

- A sequence-to-sequence (Seq2Seq) model is a type of neural network architecture that transforms an input sequence into an output sequence.
- Where both sequences can have different lengths and are often sequential data like text or speech.

# Applications of RNN in NLP

---

- **Autocompletion in Gmail:**
  - Example: When typing "not interested at," Gmail auto-completes with "this time".
  - **RNN** is used to predict and generate the next part of the sentence based on the sequence of words typed so far.
- **Language Translation (e.g., Google Translate):**
  - RNNs are widely used for translating sentences from one language to another by processing each word in a sequence.
- **Named Entity Recognition (NER):**
  - Example: In the sentence "Elon Musk is a billionaire due to Tesla's success," an RNN identifies "Elon Musk" as a person and "Tesla" as a company.
- **Sentiment Analysis:**
  - Example: A product review is analyzed to determine its sentiment (positive/negative), such as "This phone is amazing!" being classified as a positive review.

# Why Not Use Simple Neural Networks?

---

- **Challenges with Standard Neural Networks:**
  - **Sequence Information Loss:**
    - a. For instance, "How are you?" vs. "You are how?" – sequence matters in language.
  - **Variable Sentence Lengths:**
    - a. Fixed-size input layers cannot handle sentences of varying lengths effectively.
    - b. A large fixed-size input layer can lead to inefficiency or unnecessary complexity.
  - **High Computation:**
    - a. Converting words to vectors (e.g., using one-hot encoding) increases computational load.



# Sequence Modeling and Language Translation

---

- **Issue with Simple Neural Networks for Translation:**
  - If a sentence's word order is changed, the meaning changes:
    - a. "I ate pizza on Sunday" vs. "On Sunday I ate pizza".
  - **ANN Limitations:**
    - a. Cannot handle reordering or sequence-sensitive data.
  - **RNN Advantage:**
    - a. RNNs remember previous words and context, preserving sequence integrity.



# Named Entity Recognition (NER) Example

---

- **NER Explanation:**

- In the sentence: "Elon Musk loves SpaceX," an RNN identifies:
  - a. "Elon Musk" as a **person**.
  - b. "SpaceX" as a **company**.

- **How RNN Works in NER:**

- Convert words into vectors (e.g., one-hot encoding).
- Process the sentence word by word, and the RNN maintains the context of previously seen words.

# The RNN Architecture

---

- **Recurrent Structure:**

- RNNs have a loop structure where the output of each word is fed back into the network for the next word in the sequence.
- **Time Steps:**
  - a. At each step, the output carries forward the context or memory of the sequence processed so far.

- **Example:**

- For the sentence "Elon Musk loves SpaceX":
  - a. The RNN processes "Elon", "Musk", "loves", and "SpaceX" one by one, carrying context at each step.



# Deep RNN and Training Process

---

- **Training RNN for NER:**
- **Initial Weights:** Randomly initialized weights.
- **Forward Pass:** Words are processed one by one.
- **Loss Calculation:** Predicted entities are compared to actual labels (e.g., person, company).
- **Backpropagation:** Adjust weights using gradient descent to minimize the loss.
- **Epochs:**
  - Passing the entire dataset through the network multiple times to improve accuracy.



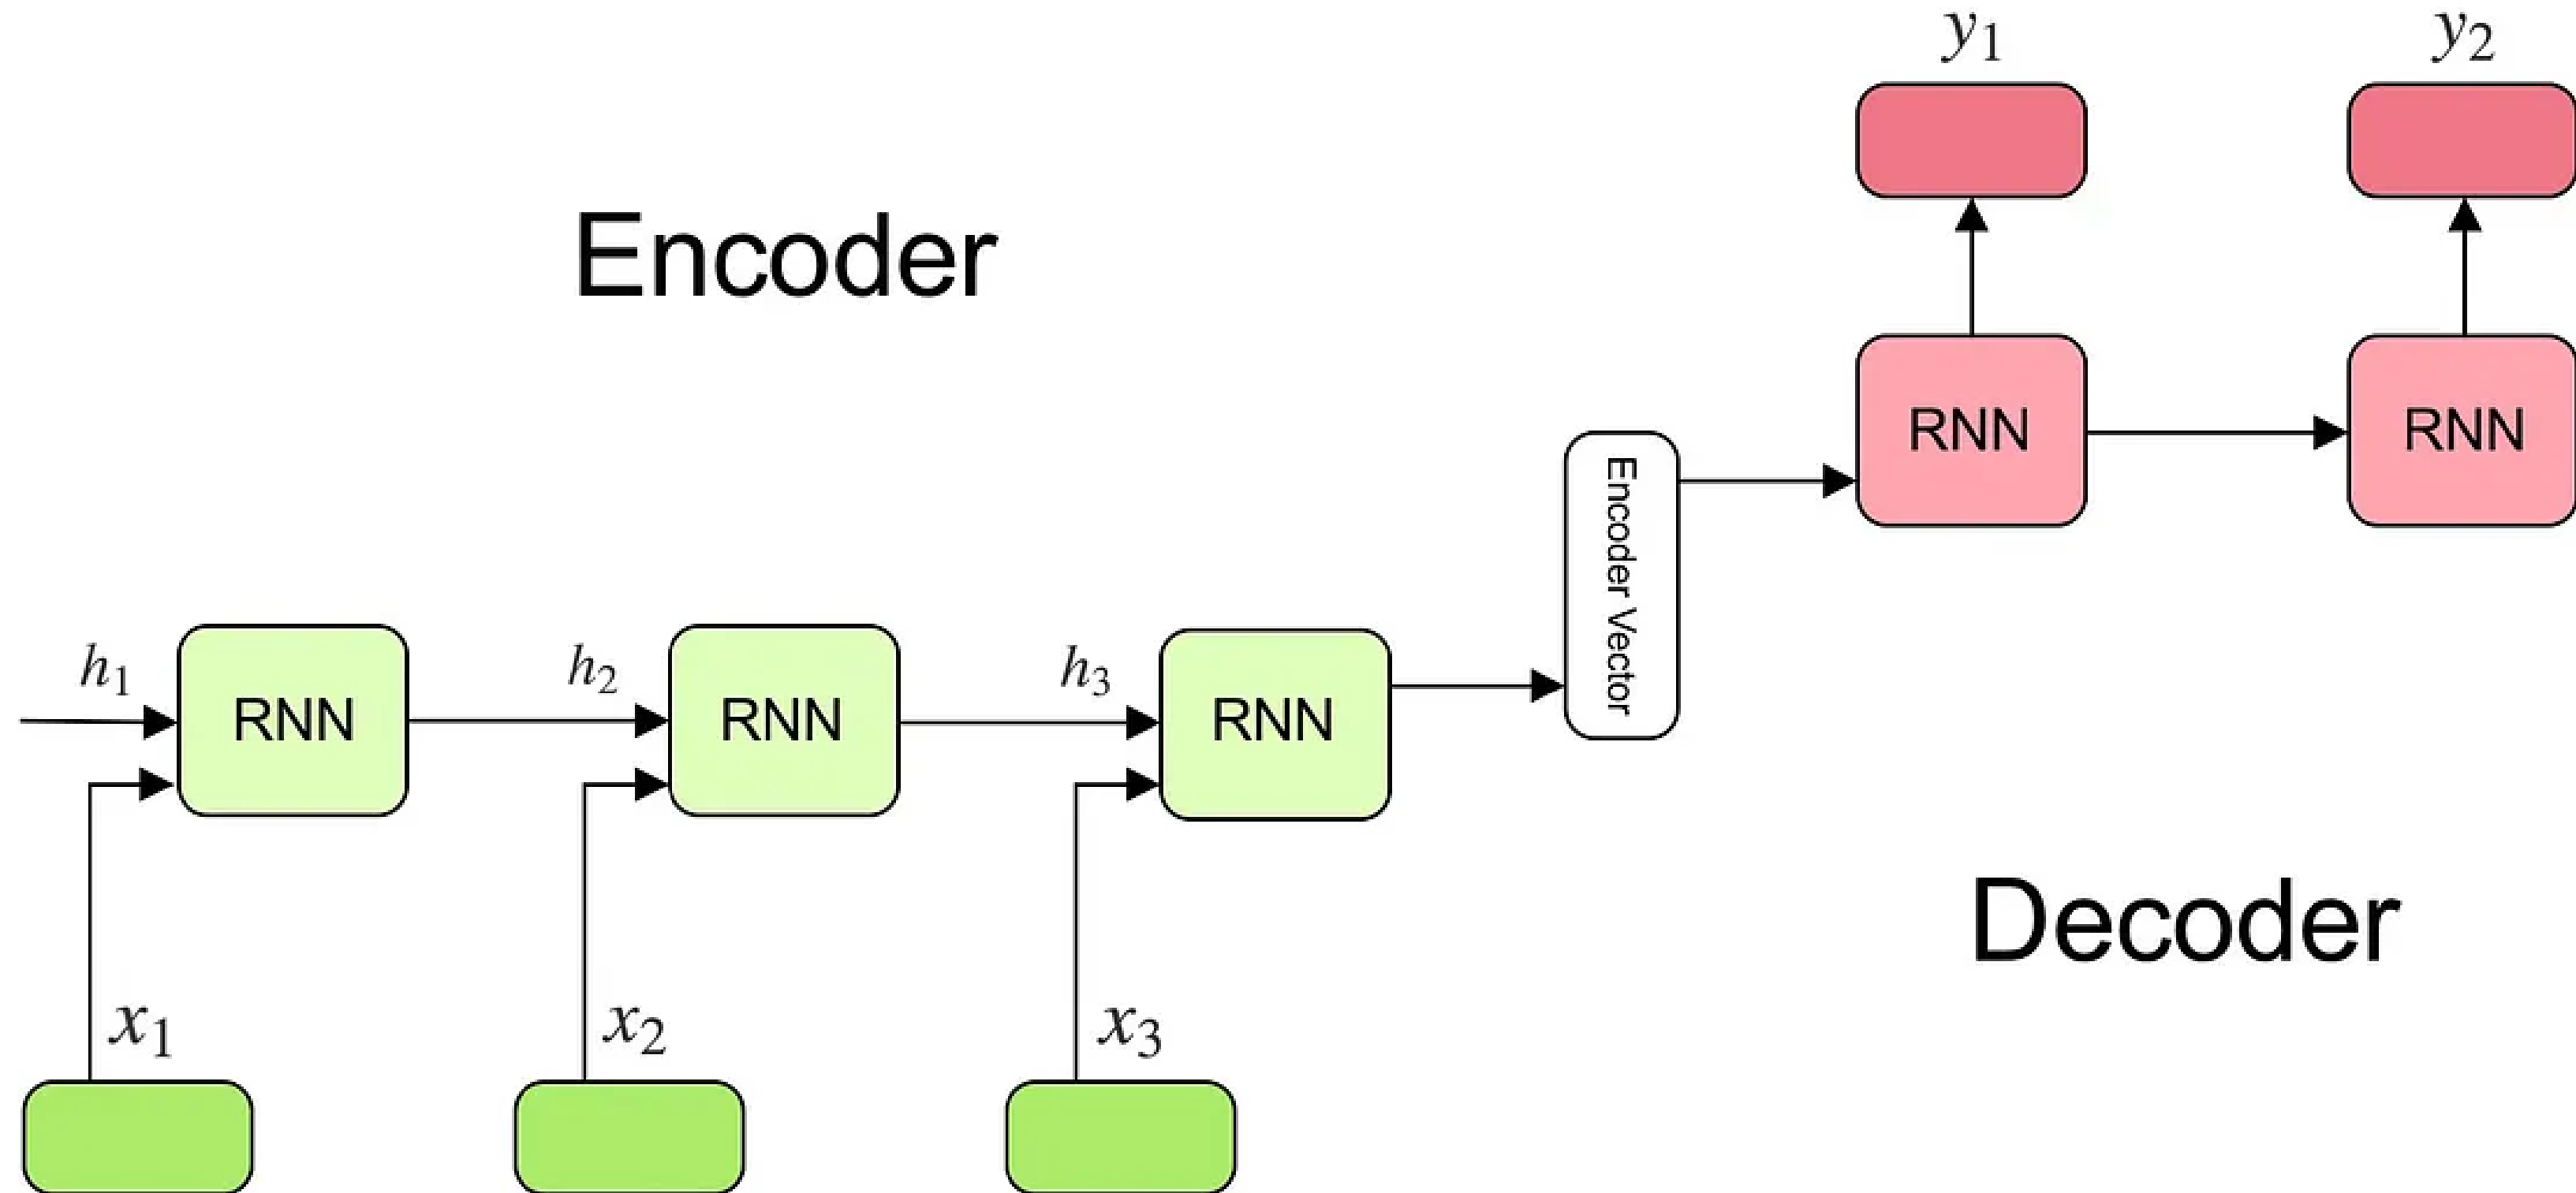
# Language Translation with RNN

---

- **Language Translation with RNN**
- **Encoder-Decoder Architecture:**
  - **Encoder:** Processes the input sequence (source language).
  - **Decoder:** Generates the output sequence (target language).
- **Process:**
  - Input sentence is fed word by word into the encoder.
  - Once the entire sentence is processed, the decoder starts translating the sentence word by word.
- **Deep RNNs:**
  - Can have multiple hidden layers to improve model performance.



# How the Sequence-to-Sequence Model works?



# Why Use RNN for Sequence Modeling?

---

- **Key Benefits of RNN:**
  - **Memory and Context:** RNNs have the ability to remember and use past information, crucial for sequence-based tasks like language translation, sentiment analysis, etc.
  - **Handling Variable-Length Sequences:** RNNs can process sequences of varying lengths, making them more flexible than traditional neural networks.

# Different Types of Recurrent Neural Networks (RNNs)

---

- Introduction to Different Types of Recurrent Neural Networks (RNNs)
- **Goal:** Understand the different types of RNN architectures used in sequence modeling tasks.
- **Many-to-Many RNN**
- **Many-to-One RNN**
- **One-to-Many RNN**
- **RNN Basics:**
  - RNNs are used for processing sequences where the order of the data matters (e.g., text, music).
  - Key feature: They can handle inputs and outputs of varying lengths.

# Many-to-Many RNN – Named Entity Recognition (NER)

- **Use Case: Named Entity Recognition (NER)**
  - **Example:** Input sentence: "Elon Musk is the CEO of SpaceX."
    - a. **Output:** Tagging of named entities: "Elon Musk" -> Person, "SpaceX" -> Organization.
- **Architecture:**
- Input sequence (e.g., "Elon Musk is the CEO of SpaceX") is processed, and each word is classified (e.g., "Elon" -> Person, "CEO" -> Role, etc.).
- This is a **Many-to-Many** RNN because each word in the input sequence has a corresponding output.
- **Generic Representation:**
  - Inputs:  $x_1, x_2, \dots, x_k$  (sequence of words).
  - Outputs:  $y_1, y_2, \dots, y_k$  (tags for each word in the sequence).
- **Other Example:**
  - **Sentence:** "Apple's stock price surged today."
  - **Entity Tags:** "Apple" -> Organization, "stock price" -> Financial Term, "surged" -> Action.

# Many-to-Many RNN – Language Translation

---

- **Use Case: Language Translation**

- **Example:** Input sentence: "I love deep learning."
  - a. **Output:** "J'adore l'apprentissage profond." (French translation).

- **Architecture:**

- The RNN processes the entire input sentence one word at a time, and after processing the last word, the model generates the translation.
- The number of words in the output can vary from the number of words in the input.
- **Generic Representation:**
  - a. Input sequence  $(x_1, x_2, \dots, x_k)$  in one language (e.g., English).
  - b. Output sequence  $(y_1, y_2, \dots, y_k)$  in the target language (e.g., French).

- **Translation Example:**

- Input: "The cat sat on the mat."
- Output: "Le chat s'est assis sur le tapis."

# Many-to-One RNN – Sentiment Analysis

---

- **Use Case: Sentiment Analysis**

- **Example:** Input paragraph: "This phone is amazing, I love it!"
  - a. **Output:** Sentiment rating: Positive (or 5 stars).

- **Architecture:**

- The RNN processes the entire paragraph (many words) and outputs a single value (the sentiment score).
- **Generic Representation:**
  - a. Input sequence  $(x_1, x_2, \dots, x_k)$  where each  $x$  represents a word.
  - b. Output ( $\hat{y}$ ): A single sentiment label (e.g., Positive, Negative, or a numerical score like 1–5 stars).

- **Other Example:**

- Input: "The movie was boring, it dragged on forever."
- Output: Negative sentiment (e.g., 1 star).

# One-to-Many RNN – Music and Poetry Generation

---

- **Use Case: Music and Poetry Generation**

- **Example:** Input seed: "Once upon a time" (for poetry generation).
  - a. **Output:** "Once upon a time, there lived a king, who ruled a vast kingdom..."
- **Example for Music:** Input seed note: A simple musical note.
  - a. **Output:** A melody created based on the initial seed note.

- **Architecture:**

- The model is fed a single input (like a seed word or note), and it generates a sequence of outputs.
- **Generic Representation:**
  - a. Input:  $x_1$  (a single word or note).
  - b. Outputs:  $y_1, y_2, \dots, y_k$  (the generated sequence of words or notes).

- **Music Generation Example:**

- Input: "C"
- Output: "C, E, G, A, C" (melody formed from the note "C").





# Summary of RNN Architectures

---

- **Many-to-Many RNN:**
  - **Input:** Sequence of words (e.g., sentence or text).
  - **Output:** Sequence of labels (e.g., tags in NER) or translated words.
  - **Examples:** NER, Language Translation.
- **Many-to-One RNN:**
  - **Input:** Sequence of words.
  - **Output:** A single value (e.g., sentiment or classification score).
  - **Examples:** Sentiment Analysis.
- **One-to-Many RNN:**
  - **Input:** A single word or note.
  - **Output:** A sequence of words or notes.
  - **Examples:** Poetry Generation, Music Composition.



# Training Process for RNN

---

- **Training:**

- **Forward Pass:** The input sequence is passed through the RNN layer by layer.
- **Loss Calculation:** Predicted output is compared with the actual output, and the loss is calculated.
- **Backpropagation:** Adjusting the weights to minimize the loss by updating the network's parameters.

- **Optimization:**

- **Gradient Descent** is used to optimize the model parameters during training.

# Training Process for RNN

---

- **Training:**

- **Forward Pass:** The input sequence is passed through the RNN layer by layer.
- **Loss Calculation:** Predicted output is compared with the actual output, and the loss is calculated.
- **Backpropagation:** Adjusting the weights to minimize the loss by updating the network's parameters.

- **Optimization:**

- **Gradient Descent** is used to optimize the model parameters during training.

# Training Process for RNN

---

- In an RNN, the **input** at each time step is denoted as  $x(t)$ , but you are also asking about the **hidden state**  $a(t)$ , which is an essential component of how RNNs process sequences. Let's break this down clearly:
- **At each time step in an RNN:**
- **Input:**  $x(t)$  is the input at time step  $t$ .
- **Hidden state:**  $a(t)$  is the hidden state at time step  $t$ , which carries the memory of the previous time steps.
- **The process of an RNN:**
- At the **first-time step** ( $t=1$ ), the RNN receives the first input,  $x(1)$ .
- The RNN starts with an initial hidden state  $a(0)$  (this is typically initialized as a vector of zeros or random values).
  - **Hidden State Update:** At time step  $t=1$ , the hidden state  $a(1)$  is computed using both the input  $x(1)$  and the previous hidden state  $a(0)$ :
- $a(1) = \text{activation function}(W \cdot [a(0), x(1)] + b)$  Where:
  - $W$  is the weight matrix.
  - $b$  is the bias term.
  - The activation function is often a **tanh** or **ReLU** function.
- So,  $a(1)$  (the hidden state at time step 1) is computed by taking the initial hidden state  $a(0)$  (which is usually zeros) and the input at time step 1,  $x(1)$ , and combining them through the RNN's weight matrix and activation function.

# RNN Unfold

