



# Natural Language Processing (NLP)

**Generative Pre-trained Transformer**

**Large Language Modeling**

By:

**Dr. Zohair Ahmed**



 [www.youtube.com/@ZohairAI](https://www.youtube.com/@ZohairAI) 

 [www.begindiscovery.com](http://www.begindiscovery.com)

# Generative Pre-trained Transformer (GPT-1)

---

- The Birth of the Decoder-Only Era Paper:
- “Improving Language Understanding by Generative Pre-Training”
- (OpenAI) Released: 11 June 2018
- First model that showed:
- Train a big Transformer to just predict the next word
- magically solves many NLP tasks with minimal fine-tuning
- **Title:** Improving Language Understanding by Generative Pre-Training
- **Authors:** Alec Radford, Karthik Narasimhan, Tim Salimans, Ilya Sutskever **Date:** June 11, 2018
- **Direct PDF link (OpenAI official):** [https://cdn.openai.com/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf](https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf)

# What is a GPT Model?


---

- GPT stands for **Generative Pre-trained Transformer**.
  - At its heart, it's a powerful text prediction engine.
  - You give it a piece of text, and its only job is to predict the most likely next word. It does this over and over to "generate" new text.
  - **Decoder-Only Architecture:** Think of the original Transformer model as having two parts: an *encoder* (for understanding input, like in translation) and a *decoder* (for generating output). GPT models are "decoder-only," meaning they are specialized solely for the task of generation.
  - **Building Blocks:** The model is built by stacking multiple identical layers, called **Transformer Blocks**.
  - Imagine a factory assembly line. Each Transformer Block is a workstation. A sentence goes into the first workstation, gets processed, and the improved version is passed to the next.
  - GPT-1 had 12 of these workstations stacked on top of each other.
- 




# What is a GPT Model?

- Now we can create a stack ( $n$ ) of modified decoder layers (called transformer block in the paper)
- Let,  $X$  denote the input sequence.
- $h_0 = X \in R^{T \times d_{\text{model}}}$
- $h_l = \text{transformer\_block}(h_{l-1}), \forall l \in [1, n]$
- Where  $h_n[i]$  is the  $i$ -th output vector in  $h_n$
- $P(x_i) = \text{softmax}(h_n[i]W_v)$




# Step 1: GPT Model

- $h_0 = X \in R^{T \times d_{\text{model}}}$
- **Start with normal words → turn them into numbers**
- The input sentence (e.g., “I am going home”) is turned into a bunch of numbers.
- We call this starting bunch of numbers  $h_0$  (h-zero). It’s just a table:
- **Rows** = words ( $T$  rows, for example 512 words max)
- **Columns** = 768 numbers that describe each word → So  $h_0$  is a big rectangle: 512 rows  $\times$  768 columns.




# Step 2: GPT Model

- We have a magic “**thinking block**” We built one smart layer (called a transformer block)
- It looks at all the words so far and makes each word’s numbers a little bit smarter.
- **We stack 12 of those magic blocks on top of each other** ( $n = 12$  in GPT-1)
- The first block takes  $h_0$  and gives us a better version  $\rightarrow h_1$
- The second block takes  $h_1$  and gives us an even better version  $\rightarrow h_2$
- ...
- The 12th block gives us the final super-smart version  $\rightarrow h_{12}$  (we call it  $h_n$ )
- That’s exactly what the equations say:
- $h_1 = \text{transformer\_block}(h_0)$
- $h_2 = \text{transformer\_block}(h_1) \dots h_{12} = \text{transformer\_block}(h_{11})$



# Step 3: GPT Model

Operation	GPU does it on the whole rectangle at once?	Time taken
<b>Make Q, K, V</b>	Yes → 3 matrix multiplications on $512 \times 768$	~0.0001 sec
<b><math>Q \times K^T \rightarrow</math> attention scores</b>	Yes → one big $512 \times 512$ matrix in one shot	~0.0002 sec
<b>Apply causal mask</b>	Yes → just subtract a $512 \times 512$ mask	instant
<b>Softmax + multiply by V</b>	Yes → back to $512 \times 768$	instant
<b>Feed-forward (<math>768 \rightarrow 3072 \rightarrow 768</math>)</b>	Yes → two big matrix multiplications	~0.0003 sec
<b>Add &amp; LayerNorm (residual)</b>	Yes → element-wise on the whole $512 \times 768$	instant




# The Training

---

- **Training Data:** GPT-1 was trained on the **Book Corpus** dataset.
- A collection of over 7,000 unique books from various genres.
- They provide long, coherent stretches of text. This is crucial for learning grammar, reasoning, and long-range dependencies (e.g., connecting a character mentioned in Chapter 1 to their actions in Chapter 5).
- It contained roughly 1 billion words. While massive at the time, this is small compared to modern models trained on trillions of words from the entire internet.
- **Tokenization - Byte Pair Encoding (BPE):** The model needs to handle any word, even ones it has never seen.
  - This allows the model to handle suffix which are frequent. Discuss Later (BERT)
- GPT-1 had a vocabulary of **40,000** such tokens.
- It Contains **12 decoder layers** (transformer blocks)
- **Context size:** 512
- **Attention heads:** 12
- **FFN hidden layer size:**  $768 \times 4 = 3072$
- **Activation:** Gaussian Error Linear Unit (**GELU**)




# The Training




# The Training


---



# The Training



# The Training



# The Training

---

- Token + Position Embeddings
- Embedding Layer Parameters
- Token Embeddings:  $|V| \times \text{embedding\_dim}$
- $40,478 \times 768 = 31 \times 10^6 = 31.1 \text{ M}$
- Position Embeddings:
  - context length  $\times$  embedding\_dim
  - $512 \times 768 = 0.3 \times 10^6 = 0.3 \text{ M}$
- Total Embedding Parameters: **31.3 M**
- *The positional embeddings are also learned, unlike the original Transformer which uses fixed sinusoidal embeddings.*
- Attention Parameters Per Block
- Attention Parameters (per transformer block)
  - Each attention head:  $W^Q = W^K = W^V$
  - $(768 \times 64): 3 \times (768 \times 64) \approx 147 \times 10^3$  parameters per head
- For 12 heads:  $12 \times 147k \approx 1.7 \text{ M}$
- Linear output layer ( $W^O$ ):  $768 \times 768 \approx 0.6 \text{ M}$
- Total per block:  $1.7 \text{ M} + 0.6 \text{ M} = 2.3 \text{ M}$
- For all 12 blocks:  $12 \times 2.3 \text{ M} = 27.6 \text{ M}$



# The Training

---

- Feed-Forward Network (FFN) Parameters
- FFN Parameters Per Block
- Two linear layers:  $2 \times (768 \times 3072) + \text{biases}$   
 $(3072 + 768) = 4.7 \times 10^6 = 4.7 \text{ M parameters per block}$
- For all 12 blocks:  $12 \times 4.7 \text{ M} = 56.4 \text{ M}$
- Final Number Everyone Remembers
- GPT-1 = ~117 Million Parameters (the first real “large” language model in 2018, tiny by 2025 standards!)

Layer	Parameters (Millions)
Embedding Layer	31.3 M
Attention Layers (12)	27.6 M
FFN Layers (12)	56.4 M
Total	116.46 M

